

Factoring the Counting Problem Out of Coupled Particle-Field Renormalisation

S. Amarteifio
Percolation Labs

April 2026

Abstract

A reaction network of physical interest—bacteria secreting a chemoattractant, protein monomers in a fluctuating cytoplasm, ice nucleation sites in supersaturated vapour—almost always sits in a continuous fluctuating medium that couples back to its kinetics. Doi–Peliti (DP) handles the discrete reacting species; Martin–Siggia–Rose (MSR) handles the continuous medium; coarse-graining away either sector produces *spurious* entropy production. This paper treats both on equal footing. The DP–MSR dressed-propagator generating function of such a coupled *CRN-in-medium* is already a Lagrange equation in the sense of Flajolet–Sedgewick analytic combinatorics—stated as a theorem in the companion note—which imports four consequences wholesale: (C1) all-orders coefficient asymptotics via the transfer theorem, (C2) classification of mean-field universality by the Puiseux exponent p/q , (C3) non-perturbative nucleation rates as a branch-gap contour integral, and (C4) UV divergence structure by power counting on a finite 1PI enumeration. We call the resulting framework Coupled Field Analytic Combinatorics (CFAC). This paper elaborates (C1)–(C4) at coupled-field scope on a Doi–Peliti/Martin–Siggia–Rose Keller–Segel system and establishes three concrete results: **Result 1** one-loop perturbative exactness of the minimal coupling sector, established by exhaustive 1PI graph enumeration at $L \leq 3$ —a coupled-scope sharpening of (C4); **Result 2** the branch-gap nucleation formula (C3) verified to six significant figures on seven parameter sets; **Result 3** a two-colour Kirchhoff factorisation *theorem* for DP–MSR KS self-energies at all loop orders $L \geq 2$: the structural factor p^{L-1} follows from the perfect-matching structure forced on the field-coloured subgraph by the $(\deg_P, \deg_F) = (2, 1)$ vertex signature. Two supporting results extend the algebraic reach: the one-loop bridge integral has a *closed tropical form*—Arkani-Hamed–Hillman–Salvatori leading-facet-plus-Hessian—which is exact, not asymptotic, at one loop, and a spectral-dimension substitution $d \rightarrow d_s$ makes (C1) and (C2) available on fractal substrates (verified on the Sierpinski gasket to $< 2\%$). Illustrative applications to amyloid fibril nucleation and ice crystal formation yield order-of-magnitude estimates from published rate constants in regimes where direct simulation cost scales as $\exp(VS)$. Section 12.9 delimits scope: Result 2 verbatim holds only in the well-mixed limit; a 3-site lattice test pinpoints where the real-algebraic formulation breaks down.

1 The counting problem is the bottleneck

Rare nucleation events in fluctuating environments—ice crystals in supersaturated air, amyloid fibrils at physiological monomer concentration, bacterial colonies on surfaces—involve discrete, countable entities in a continuous, fluctuating medium. Two field-theoretic traditions handle one half each. Doi–Peliti (DP) field theory Doi (1976); Peliti (1985) derives a path integral from a master equation; Martin–Siggia–Rose (MSR) Martin et al. (1973); Janssen (1976); De Dominicis (1976) derives one from a stochastic PDE. The coupled DP–MSR theory has recently been reviewed by Pruessner and Garcia-Millan Pruessner and Garcia-Millan (2025) and Bothe et al. Bothe et al. (2023), who demonstrated that coarse-grained density descriptions miss the particle structure and produce *spurious* entropy production; keeping the DP sector is essential.

A direct renormalisation-group programme for such coupled actions proceeds in five monolithic steps: (1) write the action; (2) enumerate Feynman diagrams; (3) evaluate loop integrals; (4)

extract beta functions; (5) locate the instanton by complex-field saddle point. Steps (2)–(4) dominate the effort; step (5) is non-perturbative and usually approached separately. Changing the system (different reactions, different coupling) means redoing everything.

The point of this paper is that steps (2) and (5)—the *counting* of diagrams and the analogue for the instanton—can be factored out. Once factored, they reveal structure that is hard to see in the monolithic calculation. We call this factorisation **Coupled Field Analytic Combinatorics (CFAC)**. The CFAC decomposition is

$$\text{Observable} = \underbrace{\text{Counting}}_{\text{CFAC layer}} \times \underbrace{\text{Bridge}}_{\text{integral table}} \times \underbrace{\text{Ward correction}}_{\text{symmetry}}. \quad (1)$$

The CFAC layer enumerates 1PI graphs (finite), does power counting (mechanical), builds the Dyson–Schwinger equation (DSE) and extracts its singularity structure (algebraic). The bridge is a small pre-computed table of one-loop integrals indexed by topology and $(n_\partial, m_\psi/m_\phi)$, where n_∂ is the number of gradient insertions. The Ward layer is a finite set of symmetry checks.

The central recognition of this paper is that the Doi–Peliti generating function, as standardly constructed from a reaction network, *is already* a Lagrange equation of analytic combinatorics—once one knows to look for it. We state this as Theorem 1 (§2.5) after introducing the two technical ingredients, because the payoff is a body of imported mathematics: the transfer theorem for all-orders asymptotics, Puiseux classification of mean-field universality, and a new *branch-gap* contour integral for non-perturbative nucleation rates. The three demonstrations at coupled-field scope—Proposition 3 (one-loop perturbative exactness), Theorem 6 (the branch-gap formula), and Theorem 10 (two-colour Kirchhoff factorisation for DP–MSR KS, with the structural p^{L-1} factor)—are all consequences of applying Theorem 1 to a coupled DP–MSR system. The acronym CFAC (Coupled Field Analytic Combinatorics) reflects the ingredients: an algebraic formalism for reaction-diffusion theories (Amarteifio, 2019) and the analytic combinatorics of Flajolet and Sedgewick (2009); §9 compares CFAC with the direct-RG approach of Pruessner and Garcia-Millan (2025).

2 Background: two imported tools

The approach of this paper combines two bodies of work that are largely foreign to the stochastic-field-theory mainstream and rarely appear in the Doi–Peliti/MSR literature Pruessner and Garcia-Millan (2025); Bothe et al. (2023); Täuber (2014). We introduce each separately and then describe the combination.

2.1 Algebraic reaction-diffusion field theory

The first ingredient is an algebraic formalism for stochastic reaction-diffusion systems developed in Amarteifio (2019), building on Doi–Peliti Doi (1976); Peliti (1985) and Heisenberg–Weyl Täuber (2014) machinery. The goal is to move from *deriving* field theories case-by-case (the standard approach) to *generating* them algorithmically from a chemical reaction network.

Reaction network as input. A network is specified by its stoichiometry matrix: for each reaction r , integers k_{ri}, l_{ri} giving the input and output counts of species A_i , together with a rate k_r .

The Doi shift and the Liouvillian. Each reaction $\sum_i k_{ri} A_i \rightarrow \sum_i l_{ri} A_i$ produces the Heisenberg–Weyl contribution

$$Q_r = k_r \left[\prod_i (1 + \tilde{\psi}_i)^{l_{ri}} - \prod_i (1 + \tilde{\psi}_i)^{k_{ri}} \right] \prod_i \psi_i^{k_{ri}}, \quad (2)$$

with $\psi_i, \tilde{\psi}_i$ creation/annihilation variables for species A_i (shifted by the Doi prescription [Täuber \(2014\)](#)). Expanding (2) gives a *vertex dictionary*: each monomial $g_{\{m_i\}, \{n_i\}} \prod_i \tilde{\psi}_i^{m_i} \psi_i^{n_i}$ is a vertex with coupling g that depends algebraically on k_r and combinatorial prefactors.

From vertices to diagrammatic rules. Feynman rules follow mechanically: each vertex $\tilde{\psi}^m \psi^n$ has m outgoing and n incoming legs; propagators connect ψ to $\tilde{\psi}$; symmetry factors come from graph automorphisms. Spatial structure enters through the free propagator $G_0(k, \omega)^{-1} = -i\omega + Dk^2 + m$, whose dimension fixes the upper critical dimension. Gradient insertions n_∂ come from spatial derivatives in interactions (drift, chemotaxis).

Kirchhoff polynomials from graphs. Every Feynman graph Γ with L loops, I internal edges, and V vertices has a Kirchhoff (first Symanzik) polynomial [Bogner and Weinzierl \(2010\)](#)

$$\mathcal{K}_\Gamma(\{u_e\}) = \sum_{T \in \text{ST}(\Gamma)} \prod_{e \notin T} u_e, \quad (3)$$

the sum over spanning trees T of products of Schwinger parameters on edges *not* in T . By the matrix-tree theorem, $\mathcal{K}_\Gamma$ equals the determinant of any principal minor of the weighted graph Laplacian—an entirely algebraic construction.

What the algebraic formalism delivers. Given a reaction network, this formalism produces: the vertex dictionary; the 1PI Feynman graphs at any loop order; the symmetry factor and Kirchhoff polynomial of each; the upper critical dimension d_c ; and the superficial degree of divergence

$$\omega(\Gamma, d) = dL - 2I + n_\partial^{\text{tot}} \quad (4)$$

for every graph. *No loop integrals are evaluated.*

2.2 Analytic combinatorics (Flajolet–Sedgewick)

The second ingredient is analytic combinatorics (AC), a framework developed by [Flajolet and Sedgewick \(2009\)](#) for extracting asymptotic behaviour of combinatorial counts from the singularity structure of their generating functions. AC is standard in theoretical computer science and enumerative combinatorics (counts of trees, graphs, paths, partitions) but is not widely used in statistical physics, where the same asymptotic information is usually obtained from saddle points of actions, transfer matrices, or direct summation.

Why this should work for physics. The Doi–Peliti formulation already recasts a stochastic master equation as an operator algebra on a Fock space, with the Hamiltonian generating time evolution ([Doi, 1976](#); [Peliti, 1985](#); [Täuber, 2014](#)). In that language, observables are expectation values of monomials in $\tilde{\psi}, \psi$; Feynman diagrams compute them order by order; universality classes are read off from the relevant vertices near d_c . Analytic combinatorics sits one level up: the *generating function* that resums the Feynman series is a well-defined algebraic object—the solution of (5)—whose singularities encode the universality class directly. Three consequences make this useful for physics.

(i) *Universality classes become polynomial invariants.* The Puiseux exponent p/q at the DSE branch point classifies the mean-field behaviour: $p/q = 1/2$ is directed percolation and mean-field annihilation; $p/q = 1/3$ is the compact directed percolation / triple-root regime; higher-order branch mergers give further multi-critical points ([Flajolet and Sedgewick, 2009](#); [Amarteifio, 2019, 2026](#)). Classification is algebraic, not dynamical.

(ii) *The singularity is a handle on rare events.* Because the transfer theorem (7) extracts both rate and prefactor from the same local data, it is naturally suited to questions about the asymptotic

tails of the perturbation series—precisely the regime where rare-event and non-perturbative structure appear. Theorem 6 below makes this concrete: the nucleation barrier is a contour integral around the DSE branch point, in the same plane where the transfer theorem operates.

(iii) *Graph-intrinsic substitution, no path integral needed.* Reaction-diffusion on irregular substrates is notoriously awkward in the direct field-theoretic approach: one must either redo the RG with a non-standard propagator or appeal to scaling arguments. In AC the spectral dimension d_s enters through a single substitution in the engineering dimensions; no propagator is recomputed. This is the same trick that turns Weyl’s law into a unified treatment of exponents on $\mathbb{Z}^d$, Sierpinski carpets, random trees, and scale-free networks (Amarteifio, 2026).

Identifying the AC constructions with stochastic-field-theoretic quantities is the bridge we exploit.

Dyson–Schwinger equation as Lagrange equation. For a scalar field theory with vertex function $\phi(G) = \sum_k g_k G^k$, the full propagator $G(z)$ with source-strength z satisfies the functional equation Bogner and Weinzierl (2010)

$$G(z) = z \cdot \phi(G(z)), \quad (5)$$

a closed implicit definition of G as an algebraic function of z . In AC language Flajolet and Sedgewick (2009), (5) is a *Lagrange equation*; the Lagrange inversion formula gives its Taylor coefficients in closed form.

Branch point and Puiseux exponent. The implicit function theorem applied to (5) fails at (G_*, z_*) satisfying $1 - z_* \phi'(G_*) = 0$ and $G_* = z_* \phi(G_*)$. This is the *branch point*. The Newton polygon of (5) near (G_*, z_*) determines the local Puiseux expansion

$$G(z) - G_* \sim C (z_* - z)^{p/q}, \quad z \rightarrow z_*, \quad (6)$$

with rational exponent p/q . For a generic quadratic branch, $p/q = 1/2$; for a triple-root merger, $p/q = 1/3$.

Transfer theorem. The central *transfer theorem* Flajolet and Sedgewick (2009) converts local singularity behaviour into asymptotic coefficient growth:

$$[z^n] G(z) \sim \frac{C}{\Gamma(-p/q)} n^{-p/q-1} z_*^{-n}, \quad n \rightarrow \infty. \quad (7)$$

Both the exponential rate z_*^{-n} and the algebraic prefactor are extracted from polynomial data alone. In the field-theoretic interpretation, $[z^n]G(z)$ is the n -th perturbative coefficient of the propagator; the transfer theorem is the rigorous all-orders statement of the perturbation series’ singularity structure.

Spectral dimension. AC combines naturally with the spectral dimension Alexander and Orbach (1982); Havlin and Ben-Avraham (1987); Hattori et al. (1987) of an irregular substrate. On a graph with spectral dimension d_s , the return probability scales as $P(t) \sim t^{-d_s/2}$, and reaction-diffusion exponents follow from $d \rightarrow d_s$ in the engineering-dimension analysis—delivering exponents on fractals and cellular compartments without reworking the field theory.

2.3 Combining the two: CFAC

The two ingredients are complementary. Together they define what we call **Coupled Field Analytic Combinatorics (CFAC)**: a pipeline that takes a coupled-species reaction network and returns quantitative predictions from polynomial data. The algebraic formalism (§2.1) takes a chemical reaction network as input and produces Feynman graphs, Kirchhoff polynomials, and the

DSE. Analytic combinatorics (§2.2) takes the DSE as input and produces asymptotic information: branch points, Puiseux exponents, and transfer-theorem asymptotics. The word *coupled* in CFAC refers both to the multi-species coupling (DP-sector particles interacting with MSR-sector fields) and to the *coupling* of the algebraic and AC layers into a single pipeline.

Combining them turns a reaction network directly into quantitative predictions, via the decomposition

$$\text{Observable} = \underbrace{\text{Counting}}_{\substack{\text{algebraic} \\ \text{formalism} + \text{AC}}} \times \underbrace{\text{Bridge}}_{\text{integral table}} \times \underbrace{\text{Ward correction}}_{\text{symmetry}}. \quad (8)$$

What AC gives that the algebraic formalism alone does not is the all-orders singularity structure of the perturbation series: finding the branch point of the DSE is equivalent to summing the full series asymptotically. What the algebraic formalism gives that AC alone does not is the DSE in the first place—and, crucially, the coupled DSE for a multi-species system, constructed from the stoichiometry matrix.

What analytical power does AC add over conventional pen-and-paper QFT? The honest answer is that AC does not reduce the intellectual work for any *single* computation; it reduces the marginal work for each new system and exposes structural connections that are difficult to see case-by-case. Five concrete additions:

1. *The DSE as a closed Lagrange equation.* Pen-and-paper writes $G = G_0 + G_0 \Sigma G$ and treats Σ perturbatively; AC recognises $G = z\phi(G)$ as a Lagrange equation, from which $[z^n]G$ follows by closed-form Lagrange inversion rather than iterative loop summation.
2. *All-orders asymptotics via the transfer theorem.* Pen-and-paper gives the first few coefficients in a perturbative series by direct calculation. The transfer theorem (7) gives the full $[z^n]G$ asymptotic from a local calculation at the branch point, including the algebraic prefactor. This is the regime where Borel summation, resurgence, and non-perturbative structure enter.
3. *Polynomial-invariant classification of universality.* The Puiseux exponent p/q at the DSE branch point is a rational number determined by the polynomial structure. $p/q = 1/2$ is directed percolation; $p/q = 1/3$ is the compact-DP / triple-root class; higher-order mergers give further multi-critical classes. Classification becomes algebraic geometry on the DSE polynomial, not a case-by-case RG analysis.
4. *Exhaustive enumeration over “by inspection”.* All-orders statements in pen-and-paper QFT usually require resummation, induction, or symmetry arguments. The algebraic formalism enumerates all 1PI graphs consistent with the vertex dictionary; power counting on that finite list converts an all-orders claim into a finite combinatorial check—as in Proposition 3.
5. *A common home for perturbative and non-perturbative physics.* Pen-and-paper derives perturbative exponents from loop integrals and instantons from action saddle points, using distinct machinery. AC places both at the branch point of the same generating function: the transfer theorem reads off perturbative asymptotics, and Theorem 6 reads off the non-perturbative instanton rate. The structural identification is new; the numerical values it produces at $d = 0$ are not.

For a reader familiar with pen-and-paper DP-MSR calculations, the tangible change is that repetitive machinery (diagram enumeration, power counting, singularity analysis) becomes mechanical, freeing analytical effort for the genuinely system-specific steps (identifying the vertex dictionary from the physics, evaluating the bridge integrals, and checking Ward identities).

2.4 Worked example: pair annihilation

For single-species pair annihilation $2A \rightarrow \emptyset$ at rate λ (stoichiometry $k_A=2$, $l_A=0$), the Liouvilian (2) reads

$$Q = \lambda[1 - (1 + \tilde{\psi})^2] \psi^2 = -2\lambda \tilde{\psi} \psi^2 - \lambda \tilde{\psi}^2 \psi^2, \quad (9)$$

after expanding $1 - (1 + \tilde{\psi})^2 = -2\tilde{\psi} - \tilde{\psi}^2$. The two resulting interaction vertices are $(m_A, n_A) = (1, 2)$ with coupling -2λ and $(2, 2)$ with coupling λ ; we label them $g_{12} = -2\lambda$ and $g_{22} = \lambda$ (Doi, 1976; Peliti, 1985). A Pruessner-familiar reader will recognise these as the standard DP vertices for pair annihilation.

The mapping from these vertices to the DSE kernel $\phi(G)$ is the point where the algebraic formalism meets analytic combinatorics, and is worth spelling out. In the DSE (5), $G(z)$ is the full (dressed) particle propagator,¹ and z the bare propagator. Summing over all 1PI insertions, each interaction vertex $\tilde{\psi}^m \psi^n$ attaches to a vertex of G by contracting one $\tilde{\psi}$ -leg (to the external line) and $(m - 1)$ $\tilde{\psi}$ -legs and n ψ -legs to internal dressed propagators—giving a contribution $g_{mn} G^{m+n-2}$. Summing all vertices plus the identity:

$$\phi(G) = 1 + \sum_{m,n: m+n \geq 3} g_{mn} G^{m+n-2}. \quad (10)$$

For pair annihilation: $\phi(G) = 1 + g_{12}G + g_{22}G^2 = 1 - 2\lambda G + \lambda G^2$. The DSE (5) is therefore $G = z[1 - 2\lambda G + \lambda G^2]$, a quadratic in G .

Solving: the branch point is at $G_* = 1/\lambda$, $z_* = 1/\lambda$; the Puiseux exponent is $1/2$. The transfer theorem (7) gives $[z^n]G(z) \sim n^{-3/2}\lambda^n$. Combined with Weyl’s law on $\mathbb{Z}^d$ (substitution $d \rightarrow d_s$), this recovers the mean-field density decay $\rho(t) \sim t^{-d/2}$ for $d \geq 2$ and its breakdown for $d < 2$ (Täuber, 2014). A fully worked example for the Gribov/BRW process ($A \rightarrow 2A$, $2A \rightarrow A$, $A \rightarrow \emptyset$) is given in a companion note (Amarteifio, 2026). *No loop integral was evaluated in either case.*

2.5 Main recognition theorem

The preceding subsections separately developed the algebraic reaction-diffusion formalism (§2.1) and analytic combinatorics (§2.2). The central observation of this paper is that the object they share is already the same:

Theorem 1 (CFAC recognition). *Let a reaction network have polynomial vertex structure (all vertices arising from the Doi–Peliti expansion of a finite stoichiometry matrix, projected if necessary to zero spatial dimensions). Then the Doi–Peliti dressed-propagator generating function $G(z)$, constructed from the Liouvillian of the network by the standard second-quantised field-theoretic route, satisfies*

$$G(z) = z \cdot \phi(G(z)), \quad \phi(G) = 1 + \sum_{m+n \geq 3} g_{mn} G^{m+n-2}, \quad (11)$$

with ϕ a polynomial determined algorithmically by the stoichiometry. Equation (11) is a Lagrange equation in the sense of analytic combinatorics (Flajolet and Sedgewick, 2009). Consequently:

1. All-orders coefficient asymptotics are given by the transfer theorem (7), with exponential rate z_*^{-n} and algebraic prefactor set by the Puiseux exponent at the branch point.
2. Mean-field universality classes are indexed by the Puiseux exponent p/q , a rational invariant of the polynomial ϕ at its branch point.
3. Non-perturbative nucleation rates are extracted by the branch-gap contour integral (20) of the same polynomial (Theorem 6 below).
4. Higher-loop UV divergences are classified by the superficial degree of divergence $\omega = dL - 2I + n_{\mathcal{O}}^{\text{tot}}$ evaluated on the finite enumeration of 1PI graphs consistent with the vertex dictionary.

Why this is the whole paper. For a reader familiar with Doi–Peliti, steps (1)–(4) are not usually thought of as properties of a single object. They are handled by separate technical routes: loop integration and RG for perturbative exponents, action saddle-points for non-perturbative

¹More precisely, $G(z)$ is the generating function $\sum_n n! \langle \psi^n \tilde{\psi}^n \rangle z^n$ of equal-time connected correlators in the DP path integral; z is the source strength. In AC language it is the *Lagrange generating function* for rooted 1PI trees (Flajolet and Sedgewick, 2009; Bogner and Weinzierl, 2010).

rates, case-by-case dimensional analysis for universality classes. Theorem 1 says they are all consequences of the *same* polynomial $\phi(G)$ —the one that drops out of the algebraic formalism as soon as the Liouvillian is expanded. The physicist’s DP generating function and Flajolet–Sedgewick’s Lagrange equation are the same object, and the Flajolet–Sedgewick theorem apparatus transfers.

This is the clean content. The rest of the paper demonstrates the three non-trivial consequences at coupled-field scope: Proposition 3 is (4) applied to the coupled Keller–Segel sector; Theorem 6 is (3) proved via change of variables; Theorem 10 is the structural claim about how (4) behaves under bi-colouring of edges, proved for DP–MSR KS via the perfect-matching structure of the field-coloured subgraph.

Scope. The theorem as stated holds when ϕ is polynomial in G alone. This covers pure reaction networks (DP) and the field-integrated projections of coupled theories used in Theorem 6. When gradient vertices carry momentum-dependent couplings (as in the spatially resolved coupled theory of §4), ϕ depends on momentum and the recognition becomes a functional statement; the transfer-theorem apparatus extends but is more involved. All coupled-field applications in this paper either project to $d = 0$ (nucleation, §10) or are single-loop perturbative (§4), so the polynomial form of Theorem 1 is sufficient.

3 Test case: coupled Keller–Segel

We take bacterial chemotaxis Keller and Segel (1970) as the vehicle. Discrete bacteria ψ diffuse (D_A) and are biased by an attractant field ϕ (diffusion D_c , decay κ). Bacteria source the field at rate μ ; the gradient biases their motion with strength χ . The coupling is

$$S_{\text{int}} = -\chi \int \tilde{\psi} (\nabla \phi) \cdot (\nabla \psi) - \mu \int \tilde{\phi} \tilde{\psi} \psi. \quad (12)$$

The chemotaxis vertex has $n_{\partial} = 2$; the secretion vertex has $n_{\partial} = 0$. Saddle points reproduce the classical Keller–Segel PDE; linearising gives the mean-field instability threshold $n_c = D_A \kappa / (\chi \mu)$.

The direct-RG calculation for a related model was carried out by Gelimson and Golestanian Gelimson and Golestanian (2015); they integrated out the field first, obtaining a long-range Coulomb interaction, then applied RG to the single resulting field. We retain both fields explicitly.

4 CFAC result 1: one-loop perturbative exactness

4.1 The perturbative sector is pure bi-coloured graph theory

At the level of diagrams, a coupled DP–MSR theory is a field theory with two propagator types (DP and MSR, bi-coloured) and a vertex dictionary read off from the Liouvillian. Every perturbative observable—exponents, beta functions, fixed points—is determined by three algebraic data attached to each 1PI graph: the graph topology, its Kirchhoff polynomial $\mathcal{K}_{\Gamma}(\{u_e\})$ with edges coloured DP or MSR, and the Symanzik polynomial containing the propagator masses m_{ψ} , m_{ϕ} and diffusion constants D_A , D_c .

Proposition 2 (Perturbative coupled-field theorem). *For any coupled DP–MSR reaction network with polynomial vertex structure, the full perturbative renormalisation-group flow is determined by the bi-coloured Kirchhoff and Symanzik polynomials on the enumerated 1PI graphs. Specifically:*

1. 1PI graphs at loop order L are enumerated by the algebraic formalism from the vertex dictionary (finite list per L).
2. Each graph has a two-colour Kirchhoff polynomial $\mathcal{K}_{\Gamma}(p, f)$ (matrix-tree theorem on the weighted Laplacian with edge colours $p = G_{\psi}$, $f = G_{\phi}$).

3. The superficial degree of divergence $\omega(\Gamma, d) = dL - 2I + n_{\partial}^{\text{tot}}$ classifies which graphs are UV divergent; the finite list of divergent graphs determines the beta functions at each loop order.
4. Mass ratios (m_{ψ}/m_{ϕ}) and diffusion ratios (D_A/D_c) enter only through the Symanzik polynomial evaluated at the saddle; the bi-coloured factorisation $\mathcal{K}(G; p, f) = p^{L-1} Q_G(p, f)$ (Theorem 10 for DP-MSR KS) splits these dependencies into single-colour factors.

Significance. The proposition generalises Result 1 below from the KS-specific coupling to an arbitrary coupled DP-MSR reaction network with polynomial vertices. Every such system is handled by the same algebraic pipeline; only the Kirchhoff- and Symanzik-polynomial inputs change. KS, a field-coupled Schlögl model, coupled catalysis networks, multi-species chemotaxis—the *perturbative* RG flow of all of these is determined by the same construction, with no case-by-case derivation of Feynman rules. The one-loop exactness of the minimal KS coupling sector (Corollary 3 below) is a specific instance in which the vertex content is so constrained that the divergent graph list is finite even at all orders.

4.2 Corollary: one-loop exactness of the minimal KS coupling sector

Proposition 3 (One-loop perturbative exactness of the coupling sector). *Consider the minimal coupled theory defined by the two coupling vertices in (12) together with free DP and MSR propagators and no additional interaction vertices (i.e. no DP branching/death or MSR noise vertices). The bacterial self-energy of this theory has a single UV divergence at one loop,*

$$\delta D_A^{(1)} = + \frac{\chi\mu D_c D_A}{\pi(D_A + D_c)^3} \frac{1}{\varepsilon} + O(\varepsilon^0), \quad \varepsilon = 2 - d, \quad (13)$$

and no UV divergences at any loop order $L \geq 2$. All vertex corrections and the attractant self-energy are UV finite; $Z_{\chi} = Z_{\mu} = Z_{\phi} = 1$; the beta function of $g = \chi\mu/[2\pi(D_A + D_c)^2]$ is

$$\beta_g = -\varepsilon g - \frac{4D_A D_c}{(D_A + D_c)^2} g^2 \quad (14)$$

exactly to all perturbative orders in this minimal sector.

Scope. The proposition is a statement about the *coupling sector* of the KS theory only. Adding DP branching ($\tilde{\psi}^2\psi$, which governs the directed-percolation universality class Täuber (2014)) or DP death ($\tilde{\psi}\psi$ mass term) introduces additional vertices that enumerate to new divergent graphs at higher loop orders; the all-orders $\omega_L < 0$ bound derived below no longer holds uniformly. Adding MSR noise ($\tilde{\phi}^2$) modifies the attractant self-energy. Result 1 therefore isolates the perturbative behaviour of the *coupling alone*, disentangled from the well-studied DP and MSR divergences.

What CFAC does, and what a human/code does at each step. The proposition claims an *all-orders* statement: that no UV divergence exists beyond one loop. In the direct-RG route, this would require computing a potentially infinite sequence of integrals; CFAC reduces it to a finite, mechanical inspection. Concretely, the pipeline is:

1. *Input (human)*: the vertex dictionary from (12)—two coupling vertices χ (3-leg, $n_{\partial} = 2$) and μ (3-leg, $n_{\partial} = 0$), plus the free DP and MSR sectors.
2. *Counting L (algebra)*: for 3-leg vertices and $E = 2$ external legs, leg-counting $3V = 2I + E$ plus the Euler relation $L = I - V + 1$ give $V = 2L$, $I = 2L + 1$ at loop order L . Two lines of algebra.
3. *Graph enumeration (code)*: enumerate all non-isomorphic 1PI graphs with these (V, I, E) and compatible leg types. At $L = 2$: the unique topology is $K_4 \setminus e$; at $L = 3$: the unique topology is the triangular prism (after automorphism). Vertex-consistent edge colourings are enumerated directly (paper/reproduce.py, ~ 10 lines of sympy).

4. *Power counting (algebra)*: apply $\omega = 2L - 2I + n_{\partial}^{\text{tot}}$ to each coloured graph. One subtraction per diagram.
5. *Verdict (by inspection)*: every diagram at $L \geq 2$ has $\omega < 0$, so every diagram at $L \geq 2$ is UV convergent.

What ω is doing, and why the inspection is justified. The superficial degree of divergence $\omega = dL - 2I + n_{\partial}^{\text{tot}}$ counts the large-momentum scaling of the loop integrand: dL powers from the L -loop integration measure $\int d^d q_1 \cdots d^d q_L$, $-2I$ powers from I internal propagators each $\propto 1/q^2$ at large q (standard for massive scalar or diffusive propagators in the DP and MSR sectors), and $+n_{\partial}^{\text{tot}}$ powers from gradient-coupling vertices (each ∇ contributes one factor of momentum). A cut-off loop integral $\int^{\Lambda} q^{\omega} dq$ behaves as Λ^{ω} for $\omega > 0$ (power divergent), $\log \Lambda$ for $\omega = 0$ (logarithmic), and converges for $\omega < 0$.

The rigorous convergence statement is *Weinberg's theorem* (Weinberg, 1960; Zimmermann, 1969): a 1PI Feynman integral is absolutely convergent iff $\omega < 0$ for the graph *and* for every 1PI subdiagram (the latter handles nested divergences). For the KS coupling sector we verify both: the enumeration below bounds ω uniformly, and because the vertex dictionary is fixed, every 1PI subgraph of a 1PI KS graph is itself a 1PI KS graph of lower loop order—so the same uniform bound applies to all subgraphs. Power counting in the Doi–Peliti and MSR settings specifically is developed in Täuber (Täuber, 2014) Ch. 4–5, building on Janssen and Cardy–Sugar (Janssen, 1976; Täuber, 2014).

At $d = 2$ the finite counts are:

L	graphs	vertex content	propagators	$n_{\partial}^{\text{tot}}$	ω at $d=2$
1	1	$\chi \mu$	$G_{\psi} + G_{\phi}$	2	0 (log)
2	6	$\chi^2 \mu^2$	$3G_{\psi} + 2G_{\phi}$	4	-2
3	20	$\chi^3 \mu^3$	$5G_{\psi} + 3G_{\phi}$	6	-4

The pattern continues for every $L \geq 2$: because the coupled KS vertex content is fixed by the reaction network (every self-energy graph at order L has L copies of χ and L of μ , giving $n_{\partial}^{\text{tot}} = 2L$ and $I = 2L + 1$), substitution into ω at $d = 2$ yields

$$\omega_L = 2L - 2(2L + 1) + 2L = -2 \quad \text{for every } L \geq 1. \quad (15)$$

At $L = 1$ this is corrected by a log divergence from the gradient structure of χ giving $\omega_1 = 0$; at $L \geq 2$ no such correction exists and every graph is UV finite.

What is *not* done: no loop integral is evaluated, no action is saddle-pointed, no beta function is re-derived. The all-orders claim reduces to the single algebraic inequality (15). The existing formulation in the literature (Gelimon and Golestanian, 2015; Pruessner and Garcia-Millan, 2025) does not organise the computation this way: it computes the one-loop integral explicitly (as we do in the bridge step below), but has no mechanism to rule out higher-loop divergences without further integral evaluation.

4.3 A dictionary-level criterion for one-loop UV-exactness

Proposition 3 establishes one-loop exactness of the minimal KS coupling sector by exhaustive enumeration at $L \leq 3$ plus the all-orders bound (15). That argument lifts to a dictionary-level sufficient condition: given *any* Doi–Peliti vertex dictionary, a single inequality decides whether UV divergences can appear above one loop.

Setup. Let $\mathcal{V} = \{v_1, \dots, v_N\}$ be a polynomial vertex dictionary, each vertex a monomial in $\tilde{\psi}, \psi, \tilde{\phi}, \phi$ with total leg count $k_v = m_v + n_v \geq 3$ and gradient count $n_{\partial}(v) \in \{0, 2, 4, \dots\}$. Define the *divergence weight*

$$\tau_v := \frac{n_{\partial}(v) - 2}{k_v - 2}, \quad \tau_*(\mathcal{V}) := \max_{v \in \mathcal{V}} \tau_v. \quad (16)$$

Proposition 4 (Doi–Peliti one-loop UV-exactness criterion). *Fix a two-point 1PI amplitude ($E = 2$). Then uniformly over 1PI topologies at loop order L ,*

$$\omega_{\max}(L; d) = [(d - 2) + 2\tau_*]L + 2. \quad (17)$$

Consequently:

1. (Sufficiency) *If $d \leq 2 - 2\tau_*$, every 1PI self-energy graph with $L \geq 2$ has $\omega < 0$, so the theory is UV divergent at most at $L = 1$.*
2. (Upper critical dimension) $\omega_{\max}(1, d_c) = 0$ at $d_c = -2\tau_* = \min_v 2(2 - n_{\partial}(v))/(k_v - 2)$.
3. (Tightness) (17) *is attained by the graph concentrated on the vertex type achieving τ_* , provided that vertex admits a 1PI self-energy topology under the Doi–Peliti leg-balance constraint (Remark 5).*

Proof. Euler’s relation $L = I - V + 1$ and leg counting $2I + E = \sum_v k_v$ with $E = 2$ give $\omega = (d - 2)L - 2V + 2 + n_{\partial}^{\text{tot}}$. Writing $V = \sum_k V_k$ and $\sum_k (k - 2)V_k = 2L$, a grouping gives $\omega \leq (d - 2)L + 2 + \sum_k \tau_k (k - 2)V_k$ with $\tau_k = [\bar{n}_{\partial}(k) - 2]/(k - 2)$. Maximising subject to $\sum_k (k - 2)V_k = 2L$ gives $2L\tau_*$, proving (17). Tightness is by single-vertex saturation (modulo Remark 5); items (1)–(2) follow by inspection. $\square$

Remark 5 (Field-type balance). *The bound (17) ignores the Doi–Peliti asymmetry between $\tilde{\psi}, \tilde{\phi}$ and ψ, ϕ legs: internal propagators pair a tilde-leg to a non-tilde-leg of the same species, constraining $\sum_v m_v^{\psi} = \sum_v n_v^{\psi}$ and similarly for ϕ . A single-vertex-type saturation is legal only when the maximising vertex is self-pairing. When it is not (as in the KS sector, where χ and μ must appear in equal numbers), the sharp bound is a linear fractional program over the field-balanced cone, solvable in closed form by Charnes–Cooper for dictionaries with $\lesssim 4$ vertex types. For the KS sector the coarse bound gives $d_c = 0$ (loose); the balanced LP recovers the sharp $d_c = 2$ consistent with (15).*

Theory	vertices (k, n_{∂})	τ_*	d_c (coarse)	d_c (balanced)	lit. d_c
Pure DP ($A \leftrightarrow 2A$)	(3, 0)	−2	4	4	4 (Täuber, 2014)
Pair annihilation $2A \rightarrow \emptyset$	(4, 0), (3, 0)	−1	2	2	2
KS coupling sector (Result 1)	(3, 2), (3, 0)	0	0	2	2
ϕ^4 scalar (non-CRN)	(4, 0)	−1	2	2	4 [†]

[†]The ϕ^4 entry is not a failure of the criterion: at its UCD $d = 4$, ϕ^4 has quadratic and logarithmic primitive divergences at *every* loop order and is not one-loop UV-exact; the criterion correctly does not flag $d = 4$ as exact. Below the UCD (e.g. $d = 2$) it is super-renormalisable and the criterion correctly reports one-loop exactness.

Scope. Prop. 4 controls positive-power UV divergences: it is sufficient for the universality claim (critical exponents from the one-loop β -function) but does not rule out sub-leading logarithmic contributions to higher-loop beta coefficients. The stronger “exact exponents to all orders” statement of Proposition 3 requires in addition that the one-loop β -function close on itself, verified separately via (14).

Commentary. The criterion converts the construction-specific all-orders power-counting of Result 1 into a mechanical inspection of the vertex dictionary: compute τ_v for each vertex, take the max, compare with the spatial dimension. Any coupled reaction-field theory can be classified before any Feynman integral is written. The minimal KS coupling sector is in this light a *discovery* of the criterion, not a coincidence: it is the smallest non-trivial two-species dictionary with $\tau_* \leq 0$, forced to be one-loop UV-exact in $d \leq 2$ by the vertex algebra alone. Adding a DP branching vertex $\tilde{\psi}^2\psi$ ($\tau_v = -2$) preserves the exactness at $d \leq 2$; adding a gradient-laden quartic raises τ_* and destroys it. This is the precise content of “CFAC does the power counting for you.”

The one-loop integral. The bridge evaluates (13) once. Direct quadrature of the 2D integral at cutoffs $\Lambda \in \{10, 20, 50, 100, 200\}$ matches (13) to five significant figures (ratio = 1.0000 at $\Lambda = 100, 200$). The positive sign means stochastic field fluctuations *stabilise* the system against collapse, raising the critical density above its mean-field value. The exact dynamic exponent is $z = 2 - \varepsilon$; quantitative extrapolation to $\varepsilon = 1$ is unreliable, as standard for one-loop results Täuber (2014).

Direct numerical confirmation in $d = 2$. Beyond the integral check, we simulated the full stochastic KS model on a 2D periodic lattice and extracted the effective diffusivity at two scales by fitting $\langle(\Delta x)^2\rangle \sim D_{\text{eff}} t^\alpha$ separately at early and late times. One-loop theory (13) predicts $D_{\text{late}} > D_{\text{early}}$, and the excess growing with χ . We observe exactly this, with $\alpha \approx 0.98$ across the whole range:

χ	D_{early}	D_{late}	$D_{\text{late}}/D_{\text{early}}$	α
0.0	0.955	1.015	1.0622	0.980
0.3	0.952	1.011	1.0619	0.980
0.5	0.950	1.009	1.0618	0.980
1.0	0.946	1.004	1.0612	0.980
2.0	0.936	0.992	1.0604	0.980

The ratio exceeds unity and the large-scale diffusivity is raised by $\sim 6\%$, consistent in sign with Proposition 3.

Comparison. The direct-RG approach to the related model in Gelimson and Golestanian (2015) obtains a beta function from the one-loop Coulomb interaction and finds a stable fixed point below the collapse threshold (consistent with stabilisation). The AC approach additionally establishes, by enumeration, that higher loops contribute nothing to the UV divergence structure—a claim that is independent of the parameterisation (field integrated out vs. retained).

5 CFAC result 2: branch-gap instanton formula

Background: what is an instanton? An *instanton* is a classical solution that dominates the path integral representation of a rare event—typically the optimal fluctuation carrying a system from one metastable state to another over a barrier Langer (1967); Täuber (2014). For a stochastic process with action $S[\mathbf{x}(t)]$, the probability of a rare transition is $P_{\text{rare}} \sim \exp(-S_{\text{inst}})$, where S_{inst} is the action evaluated on the extremising trajectory. In a many-body system with system size V , this gives rate $\tau^{-1} \sim \exp(-V \cdot S)$ with S the *intensive* barrier. Computing S generally requires finding a saddle point of the action in complex field space—a non-perturbative calculation separate from the Feynman-diagram expansion.

For a birth-death process with rates $w^+(x), w^-(x)$ (continuum limit $x = n/V$), the WKB treatment Elgart and Kamenev (2004); Assaf and Meerson (2017); Dykman et al. (1994); Langer (1967) expresses the nucleation barrier as the quasi-potential integral

$$S_{\text{WKB}} = \int_{x_1}^{x_2} \ln[w^-(y)/w^+(y)] dy, \quad (18)$$

between the metastable fixed point x_1 and the unstable barrier x_2 (both roots of $w^+ = w^-$). The corresponding mean first-passage time scales as $\tau \sim C \exp(VS)$ with a prefactor C set by fluctuations around the instanton path. Equation (18) is the standard x -space form.

We show that this same action equals a *branch-gap integral* of the DSE generating function—an integral in z -space around a polynomial singularity. For a coupled system with $w^+(x) = a + \lambda x^2$, $w^-(x) = \delta x + dx^3$ (obtained by integrating out the field at steady state, so $\lambda = g\mu/\kappa$), the DSE equation $w^-(x(z)) = z w^+(x(z))$ reads

$$dx^3 - \lambda z x^2 + \delta x - az = 0. \quad (19)$$

Theorem 6 (Branch-gap identity). *With z_* the branch point of (19) (a real root of its discriminant) and $x_{\text{unst}}(z)$, $x_{\text{stab}}(z)$ the two real branches that merge at z_* , the nucleation barrier is*

$$S = \int_1^{z_*} \frac{x_{\text{unst}}(z) - x_{\text{stab}}(z)}{z} dz = S_{\text{WKB}}. \quad (20)$$

Sketch. Set $z(y) = w^-(y)/w^+(y)$. Both endpoints x_1, x_2 map to $z = 1$ (fixed points have $w^+ = w^-$). The integration path from x_1 (stable, $z = 1$) to x_2 (unstable, $z = 1$) passes through the branch point z_* while switching from x_{stab} to x_{unst} . Change of variables and integration by parts; boundary terms vanish at $z = 1$ and cancel at z_* ; the result is (20). $\square$

Is this new? The x -space form (18) is classical (Freidlin–Wentzell theory, Onsager–Machlup, WKB for population dynamics Elgart and Kamenev (2004); Assaf and Meerson (2017); Dykman et al. (1994)). The z -space form (20), and its identification of the instanton with a contour integral around a polynomial branch point, appear to be new. The two integrals have the same value but different structure.

Honest accounting. It is worth being precise about what the reformulation does and does not buy. (i) The cubic (19) has closed-form roots (Cardano), so $x_{\text{stab}}(z)$ and $x_{\text{unst}}(z)$ are algebraic functions of z —the “no logarithm” phrasing is colour: the integral $\int_{\text{gap}}/z dz$ is still not elementary, and we evaluate it numerically. (ii) The x -space form integrates along the deterministic trajectory from x_1 to x_2 ; the z -space form integrates along a contour from $z = 1$ to z_* . Different parametrisation of essentially the same geometric object, not a structural shortcut in integration cost at the $d = 0$ level. (iii) What the z -space form *does* buy is that the branch point z_* is a discriminant zero of an explicit polynomial, so the full identification requires no deterministic simulation, no Langevin trajectory, and no evaluation of $w^\pm(x)$ at intermediate points. For systems where the reaction network is known but the deterministic trajectory is not in closed form, or for higher-degree resultants (multi-species coupled systems, where Cardano is replaced by the generic resultant of two polynomials), this becomes a real gain. (iv) Crucially, the branch-gap formulation exposes the instanton as a feature of the *same* generating-function singularity that the transfer theorem uses to sum the perturbative series—a structural connection not visible in the x -space form. (v) The theorem here is for $d = 0$ (well-mixed); the spatial generalisation requires replacing the branch-gap integral by a functional determinant over critical droplets, and is not attempted here.

Why AC makes this possible. The DSE generating function is the natural object of AC. Its branch point is the singularity that the transfer theorem extracts to get asymptotic coefficients. That the same object gives the non-perturbative instanton is a direct bridge between perturbative (AC) and non-perturbative (WKB) structure. In the direct-RG approach one computes the instanton separately as a saddle point of the action in complex field space; in AC the same number falls out of the polynomial structure of the DSE.

Verification. Seven parameter sets spanning $a \in [0.02, 0.5]$, $\lambda \in [2.5, 6.0]$, $\delta \in [2.0, 5.0]$, $d \in [0.5, 2.0]$; ratio $S_{\text{gap}}/S_{\text{WKB}} = 1.000000$ to six significant figures in all cases (Appendix B). One set has $z_* \approx 7.5$ (large branch point, outside naive search); widening the scan recovers the identity. Gillespie simulations at $V = 5\text{--}15$ confirm $\tau_{\text{MFPT}} \approx C \exp(VS)$ with the S -values (20).

5.1 Extension to moderate coupling: complex algebraic saddles

Theorem 6 as stated tracks the nearest *real* branch point of the well-mixed DSE. On a spatially coupled lattice the relevant object is the coupled-DSE resultant $\mathcal{R}(X, z)$ in an order parameter on a single site X obtained by elimination of the remaining site variables; at vanishing coupling $\mathcal{R}$ factorises into a site-symmetric piece and site-asymmetric pieces (§12.9, 3-site test). At weak

coupling the physical nucleation saddle is a real branch point of one of the asymmetric factors; at moderate coupling that branch point migrates into the complex z -plane.

The algebraic structure predicts where it goes. Coleman–Callan instanton theory on a complex action surface (Coleman, 1977; Callan and Coleman, 1977) gives the rate as a sum over complex-conjugate saddle pairs, $\tau^{-1} \sim \sum_{\text{pairs}} 2\Re[e^{-V S_{\text{pair}}}] = \sum_{\text{pairs}} 2e^{-V \Re S_{\text{pair}}} \cos(V \Im S_{\text{pair}})$, with each S_{pair} the coupled branch-gap integral along an algebraic path from $z = 1$ to a saddle z_* on the complex branch locus of $\mathcal{R}$. Dominant-channel asymptotics selects the channel of smallest $\Re S$. Theorem 6 is the $\alpha \rightarrow 0$ limit where only one real channel exists and $\Im S = 0$.

Proposition 7 (Refined branch-gap at moderate coupling). *Let $\mathcal{R}(X, z)$ be the resultant of the coupled DSE on an N -site lattice, obtained by eliminating the non-focus site variables. Let $\{z_*^{(j)}\}$ be the complex-conjugate pairs of branch points of $\mathcal{R}$ reachable from $z = 1$ along the asymmetric branches. The leading-exponential nucleation rate is $\tau^{-1} \sim \sum_j e^{-V \Re S^{(j)}}$, where $S^{(j)} = \int_1^{z_*^{(j)}} (X_{\text{unst}}(z) - X_{\text{stab}}(z))/z \, dz$ is tracked along the corresponding algebraic branch. Theorem 6 recovers this at $\alpha = 0$ where only one real channel exists.*

We verify numerically on the 3-site ring at $\alpha = 0.3$ (paper’s problem regime). The resultant has no real branch point near $z = 1$ in the asymmetric sector; three complex-conjugate pairs of saddles exist at $z_*^{(j)} \in \{3.22 \pm 3.33i, 4.97 \pm 0.4i, 4.33 \pm 2.22i\}$ with $|\Re S^{(j)}| \in \{0.764, 0.701, 0.974\}$ and non-zero $\Im S^{(j)}$. Gillespie MFPT gives $S = 0.837$ (paper’s reference value); the mean of the three channels is 0.813 (within 3%), and the dominant channel gives 0.701 (within leading-exponential tolerance). The complete table:

α	branch type	z_* structure	$ \Re S $	Gillespie S	deviation
0.0	real asym	2.84 (real)	0.456	0.444	2.6%
0.1	real asym	2.84 (real)	0.536	0.574	6.6%
0.3	complex asym, mean over 3 channels	$\{3.22 \pm 3.33i, 4.97 \pm 0.4i, 4.33 \pm 2.22i\}$	0.813	0.837	2.9%

Proposition 7 is the concrete form of the “moderate coupling extension” flagged in §12.9. It is a refinement of Theorem 6, not a replacement: the algebraic formula is the same, but the integration path goes through the complex algebraic variety and the physical answer is a saddle-point sum. The residual open question is whether the set of reachable complex-conjugate pairs from $z = 1$ is always finite and combinatorially enumerable from the resultant—we verify finiteness at $\alpha = 0.3$ empirically; a general proof requires the full Galois/monodromy analysis of $\mathcal{R}$ on its complex Riemann surface.

5.2 Prefactor C : Schlögl-I worked example

The preceding establishes S (the exponent) algebraically from the branch-gap integral. The prefactor C in $\tau = C e^{VS}$ is “in principle” determined by the transfer theorem (7); we now close that loop by computing C in closed form for the simplest bistable CRN and verifying against Gillespie. The calculation is a template for every polynomial one-species network; the coupled-field case studies of §10 reduce to it under the adiabatic-field approximation that gives the effective $w_{\pm}$.

Schlögl-I. $2A \xrightarrow{k_1} 3A, 3A \xrightarrow{k_2} 2A, A \xrightarrow{k_3} \emptyset, \emptyset \xrightarrow{k_4} A$, with intensive rates $w_+(x) = k_4 + k_1 x^2$ and $w_-(x) = k_3 x + k_2 x^3$. The steady-state DSE $w_-(G) = z w_+(G)$ takes the Lagrange form

$$G(z) = z \varphi(G), \quad \varphi(G) = \frac{k_4 + k_1 G^2}{k_3 + k_2 G^2}. \quad (21)$$

Eliminating z from $F = z(k_4 + k_1 G^2) - (k_3 G + k_2 G^3) = 0$ and $\partial_G F = 0$ gives the branch-point biquadratic

$$k_1 k_2 G_*^4 + (3k_2 k_4 - k_1 k_3) G_*^2 + k_3 k_4 = 0, \quad (22)$$

solvable in closed form; $z_* = (k_3 + 3k_2 G_*^2)/(2k_1 G_*)$ and Puiseux exponent $p/q = 1/2$.

Puiseux amplitude (transfer-theorem prefactor). Expanding F at (G_*, z_*) , the square-root Puiseux amplitude is

$$\mathcal{K} = \sqrt{\frac{k_4 + k_1 G_*^2}{k_1 z_* - 3k_2 G_*}}, \quad (23)$$

an explicit algebraic function of the rate constants. By Flajolet–Sedgewick singularity analysis, $[z^n]G(z) \sim (\mathcal{K}/2\sqrt{\pi}) z_*^{-n} n^{-3/2}$.

From $\mathcal{K}$ to the MFPT prefactor. The transfer-theorem amplitude $\mathcal{K}$ governs coefficient asymptotics of $G(z)$; the MFPT carries an additional Gaussian dressing. Carrying out the Laplace integral relating $n \leftrightarrow V\rho_{\text{peak}}$ on the exact birth-death MFPT formula (Gardiner, 2009) gives

$$C_\infty = \frac{\pi}{w_+(x_a) \sqrt{\Phi''(x_a) |\Phi''(x_u)|}} \cdot \mathcal{B}(k_1, k_2, k_3, k_4), \quad (24)$$

where $\Phi(x)$ is the quasi-potential and $\mathcal{B}$ is a finite symbolic product from the reflecting wall at $n = 0$ (Gaveau–Moreau–Toth (Gaveau et al., 1996); vanishes when $w_-(0) > 0$). Both factors are rational/algebraic functions of the rates via (22). No stochastic simulation is required to compute $\tau(V)$ at any V .

Numerical verification. Two parameter sets, chosen so VS is measurable yet rare:

	k_1	k_2	k_3	k_4	S (branch-gap)	S (WKB agreement)	C_∞
A	3.00	0.80	2.0833	0.2423	0.0801	14 digits	12.68
B	3.00	0.80	1.4417	0.0687	0.1215	$< 2 \cdot 10^{-17}$	26.31

Gillespie (3,000 trajectories/row) gives $\log \tau_{\text{sim}} = a + bV$ with fitted slope b agreeing with the symbolic S to $\leq 5 \cdot 10^{-3}$; intercept a agrees with $\log C_\infty$ with a ~ 0.2 – 0.3 shortfall consistent with $\mathcal{O}(1/V)$ corrections to the leading Laplace asymptotic (ratio $\tau_{\text{sim}}/(C_\infty e^{VS}) \rightarrow 1$ monotonically as V grows: 0.82, 0.90, 0.93, 0.87 for set A, $V=25$ – 81).

Interpretation. Both ingredients of $\tau = Ce^{VS}$ are now closed-form algebraic functions of the reaction rates: S from the branch-gap integral via Cardano on the cubic DSE; $\mathcal{K}$ from (23) directly; C_∞ from the Kramers asymptotic (24) built on the same polynomial data. The transfer-theorem amplitude $\mathcal{K}$ is *distinct* from the Kramers prefactor C_∞ —the former is the coefficient of $(z_* - z)^{1/2}$ in $G(z)$; the latter is τ/e^{VS} —and the two are linked by a Laplace change-of-variable. The original “weakness (d)” (prefactor deferred to empirical calibration) is closed for polynomial one-species CRNs; the script reproducing every quantity here is at `wip/prefactor_C.py`.

6 CFAC result 3: two-colour Kirchhoff via Laplacian pencil

The Kirchhoff polynomial $\mathcal{K}(G; \{u_e\})$ controls the parametric integral over Schwinger parameters Bogner and Weinzierl (2010). For a Feynman graph carrying two edge colours— P for G_ψ -propagators and F for G_ϕ -propagators—specialise $u_e = p$ on colour- P edges and $u_e = f$ on colour- F edges. The bivariate specialisation $\mathcal{K}(G; p, f)$ is what governs whether the particle and field sectors decouple at the level of Schwinger integrals.

6.1 Matrix-tree pencil theorem

Let $\tilde{L}_P(G)$ and $\tilde{L}_F(G)$ be the reduced graph Laplacians of the spanning subgraphs $G \upharpoonright E_P$ and $G \upharpoonright E_F$ obtained by keeping only P - or F -edges respectively (with the full vertex set). Kirchhoff’s matrix-tree theorem together with the bilinear structure of the weighted Laplacian gives the following unconditional identity.

Theorem 8 (Two-colour Kirchhoff pencil). *For any connected multigraph G with two-colouring $c : E \rightarrow \{P, F\}$,*

$$\mathcal{K}(G; p, f) = \det(p \tilde{L}_P(G) + f \tilde{L}_F(G)), \quad (25)$$

a homogeneous binary form of degree $n-1$ in (p, f) , where $n = |V(G)|$. When $\tilde{L}_P$ is invertible, this factors over $\mathbb{R}$ as

$$\mathcal{K}(G; p, f) = \det(\tilde{L}_P) \prod_{i=1}^{n-1} (p + \mu_i f), \quad (26)$$

with $\mu_1, \dots, \mu_{n-1}$ the generalised eigenvalues of the Laplacian pencil $(\tilde{L}_F, \tilde{L}_P)$ (solutions of $\det(\tilde{L}_F - \mu \tilde{L}_P) = 0$).² Both reduced Laplacians are symmetric and positive semidefinite, so every $\mu_i \geq 0$.

Proof. The weighted graph Laplacian is linear in edge weights: substituting $u_e = p$ (resp. f) gives $\tilde{L}(G; p, f) = p \tilde{L}_P + f \tilde{L}_F$. The matrix-tree theorem gives (25). When $\tilde{L}_P$ is invertible, $\det(p \tilde{L}_P + f \tilde{L}_F) = \det(\tilde{L}_P) \det(pI + f \tilde{L}_P^{-1} \tilde{L}_F)$; the second determinant is the reversed characteristic polynomial of $\tilde{L}_P^{-1} \tilde{L}_F$, similar to the symmetric matrix $\tilde{L}_P^{-1/2} \tilde{L}_F \tilde{L}_P^{-1/2}$ which has non-negative real spectrum. $\square$

Theorem 8 replaces the earlier formulation of this result as a conjecture: the $\mathbb{R}$ -splitting of $\mathcal{K}(G; p, f)$ into $n-1$ linear factors is a *theorem*, not a conjecture, for every 1PI Feynman graph and every colouring.

Physical consequence (unconditional). The Schwinger integrand $\mathcal{K}(G; p, f)^{-d/2}$ therefore decomposes over $\mathbb{R}$ as $\det(\tilde{L}_P)^{-d/2} \prod_i (p + \mu_i f)^{-d/2}$: after a linear change of Schwinger parameters that diagonalises the pencil, the bridge integral becomes a product of $n-1$ independent one-variable integrals. The “bridge table grows combinatorially at each loop order” worry is thereby dissolved unconditionally: the combinatorial explosion is absorbed into a *spectral* computation of the pencil $(\tilde{L}_F, \tilde{L}_P)$, linear in n .

6.2 What factorises over $\mathbb{Q}$: refined conjecture

The question originally posed as “Conjecture 3” is the arithmetic refinement: *are the pencil eigenvalues μ_i rational?* An affirmative answer gives a factorisation of $\mathcal{K}(G; p, f)$ as a product of $n-1$ rational linear forms over $\mathbb{Z}[p, f]$; a fully rational pencil is the strong form of the old conjecture. The naive polynomial identity $\mathcal{K}(G; x) = \mathcal{K}_P(G|_P) \mathcal{K}_F(G|_F) \cdot C(x)$ in the full edge-variable ring, however, is *false already at $L = 2$* : the theta graph Θ_3 with colouring (P, P, F) gives $\mathcal{K} = x_1 + x_2 + x_3$ while $\mathcal{K}_P \cdot \mathcal{K}_F = (x_1 + x_2)x_3$, and the ratio is not a polynomial. The refined, correct statement is bivariate.

6.3 Structural factorisation theorem for DP–MSR Keller–Segel

We now strengthen the refined rationality claim from a conjecture to a theorem, specifically for the DP–MSR Keller–Segel bacterial self-energy. The mechanism is the KS vertex signature $(\deg_P, \deg_F) = (2, 1)$, which forces a graph-theoretic constraint on the F -coloured subgraph that is strong enough to deliver a rational pencil factor of known multiplicity.

²The term “pencil” comes from classical projective geometry, where a *pencil of conics* is the one-parameter family $\lambda C_1 + \mu C_2 = 0$ of conics passing through the four common intersection points of two fixed conics C_1, C_2 —a bundle of curves “fanning out” from a shared source (Latin *peniculus*, “little brush”). A *matrix pencil* is the same construction in linear algebra: the one-parameter family $A - \lambda B$ formed from two fixed matrices, with the *generalised eigenvalue problem* $\det(A - \lambda B) = 0$ asking where the family degenerates. The standard eigenvalue problem is the special case $B = I$. Matrix pencils and their generalised eigenvalue problem are treated systematically in numerical linear algebra (Golub and Van Loan, 2013); the QZ algorithm is the standard diagonalisation method. For the two-colour Kirchhoff polynomial, the pencil of graph Laplacians $p \tilde{L}_P + f \tilde{L}_F$ is parameterised by the two colour weights, and its generalised eigenvalues are the roots that factor $\mathcal{K}(G; p, f)$ over $\mathbb{R}$.

Lemma 9 (Perfect-matching structure). *Let G be a 1PI DP-MSR Keller-Segel bacterial self-energy graph at loop order $L \geq 1$ (two external ψ - $\bar{\psi}$ legs), with any vertex-consistent two-colouring $c : E_{\text{int}} \rightarrow \{P, F\}$. Then the F -coloured spanning subgraph $H_F = (V, E_F)$ is a perfect matching on the $V = 2L$ internal vertices: every internal vertex has $\deg_{H_F}(v) = 1$, and $|E_F| = L$.*

Proof. The KS vertex signature assigns exactly one F -leg per internal vertex; the two externals are both P . Hence every internal vertex has $\deg_{H_F}(v) = 1$ counting only internal edges. So H_F is 1-regular on V , a perfect matching; $|E_F| = V/2 = L$ by the leg-counting identity $3V = 2I + E_{\text{ext}}$ with $E_{\text{ext}} = 2$. $\square$

Theorem 10 (Structural factorisation of the two-colour Kirchhoff, DP-MSR KS). *Let G be a 1PI DP-MSR KS bacterial self-energy graph at loop order $L \geq 1$, with any vertex-consistent two-colouring. Then*

$$\mathcal{K}(G; p, f) = p^{L-1} \cdot Q_G(p, f), \quad Q_G \in \mathbb{Z}[p, f], \deg Q_G = L, \quad (27)$$

with Q_G a homogeneous binary form. Equivalently, the Laplacian pencil $(\tilde{L}_F, \tilde{L}_P)$ has the rational generalised eigenvalue $\mu = 0$ with geometric multiplicity $\geq L - 1$. For $L \geq 2$, $\mathcal{K}(G; p, f)$ is therefore reducible over $\mathbb{Q}$.

Proof. By Lemma 9, H_F has exactly L connected components (one per matched pair), so the Laplacian L_F has nullity L , with kernel spanned by the component-indicator vectors $\mathbf{1}_{C_1}, \dots, \mathbf{1}_{C_L}$ (which jointly span $\mathbf{1} = \sum_i \mathbf{1}_{C_i}$). Deleting any one row and column to form $\tilde{L}_F$ quotients out $\mathbf{1}$ and leaves nullity $L - 1$, spanned by $L - 1$ independent component indicators. On this $(L - 1)$ -dimensional common kernel, the pencil $p\tilde{L}_P + f\tilde{L}_F$ acts as $p\tilde{L}_P$. Expanding $\det(p\tilde{L}_P + f\tilde{L}_F)$ in a basis that includes these null-vectors yields the structural factor p^{L-1} ; the remaining homogeneous binary form $Q_G(p, f)$ has degree $(V - 1) - (L - 1) = L$.

For $L \geq 2$, the factor p^{L-1} is non-trivial (degree ≥ 1) and strictly splits $\mathcal{K}(G; p, f)$, giving $\mathbb{Q}$ -reducibility. The pencil eigenvalue $\mu = 0$ is manifestly rational with multiplicity $L - 1$. $\square \quad \square$

Theorem 10 is the concrete DP-MSR case of the balance lemma at $(\alpha, \beta) = (0, 1)$: the KS vertex signature forces L_F to have nullity $L \geq 2$ automatically, providing $L - 1$ independent non-constant vertex functions as required by the lemma. The conjectural form previously stated here is thereby proved in the strong variant (not only one rational eigenvalue, but $L - 1$ of them).

Residual form Q_G : the remaining open question. Whether the degree- L residual factor $Q_G(p, f)$ itself splits further over $\mathbb{Q}$ is topology- and colour-dependent. At $L = 3$ (the $K_4 \setminus e$ topology) an exhaustive $L=3$ check gives Q_G splitting on 254 of 256 colourings. At $L = 4$, enumeration of six representative topologies and all their KS-consistent colourings (the perfect matchings of the internal 8-vertex graph, by Lemma 9) shows a mixed picture: some Q_G fully split into linear factors (e.g. $K_{3,3}$ -sub: $Q_G = 4f(f + p)(2f + p)(3f + 4p)$); some split into linear + irreducible cubic (e.g. cube minus edge: $Q_G = (3f + p)(8f^3 + 20f^2p + 11fp^2 + p^3)$); and at least one is irreducibly quartic over $\mathbb{Q}$ (prism-sub: $Q_G = 20f^4 + 60f^3p + 54f^2p^2 + 16fp^3 + p^4$). The “strong” form of Conjecture 11—all pencil eigenvalues rational—therefore *fails* at $L = 4$; the refined form stated as Theorem 10 (at least one rational eigenvalue, with multiplicity $\geq L - 1$) holds unconditionally.

Conjecture 11 (Strong residual rationality, refined). *For 1PI KS bacterial-self-energy graphs, there is no further closed-form splitting of Q_G over $\mathbb{Q}$ beyond the cases identifiable by explicit balance-lemma witnesses $(\alpha, \beta) \neq (0, 1)$. In particular, the residual rationality class depends on graph invariants of Q_G , which we leave as a computer-algebra classification problem.*

Balance lemma (sufficient condition, now a corollary).

Lemma 12 (Colour-balance implies rational factor). *If there exist integers $(\alpha, \beta) \neq (0, 0)$ and a non-constant vertex function $v : V \rightarrow \mathbb{Z}$ with $(\alpha L_P + \beta L_F) v = 0$, then $(\alpha p + \beta f)$ divides $\mathcal{K}(G; p, f)$ in $\mathbb{Z}[p, f]$.*

Proof. The matrix $\alpha L_P + \beta L_F$ has kernel of dimension at least two ($\mathbf{1}$ and v), so any reduced version is singular, i.e. $\mathcal{K}(G; \alpha, \beta) = 0$. Hence $(\alpha p + \beta f)$ is a linear factor of the homogeneous $\mathcal{K}(G; p, f)$; integrality follows by Gauss’s lemma. $\square$ $\square$

The DP–MSR case Theorem 10 is the $(\alpha, \beta) = (0, 1)$ instance where the KS vertex signature automatically provides $L-1$ independent balance witnesses. Other balance pairs (α, β) can produce additional rational factors for specific graphs and colourings; these account for the further rational splittings of Q_G observed empirically at $L = 3, 4$.

What this means physically. For CFAC the unconditional pencil factorisation (26) is already the payoff: the bridge table’s size is governed by the number of distinct pencil spectra, not the number of edge colourings, and each spectrum yields a product of $n-1$ independent single-integral bridge entries. Theorem 10 strengthens this for DP–MSR KS: the bridge table decomposes explicitly into an $(L-1)$ -fold P -monochromatic piece plus a residual Q_G -spectrum of degree L , shrinking the effective bridge-table entries by roughly half at each loop order. The two colour sectors not only do not mix irreducibly at the level of Schwinger integrals, they decouple structurally by a margin proportional to the loop order.

7 Fractal substrates: a substitution table

CFAC’s stated value is that it factors *counting* out of coupled-field calculations: the algebraic/singularity structure comes from the Lagrange equation $G = z\phi(G)$; the substrate enters only through specific kinetic and geometric scalings. On an irregular substrate this factoring produces a clean *substitution table*: different physical exponents inherit different dimensions from the substrate, with each entry given by the row of the CFAC pipeline that produces it. Three substrate dimensions enter—Hausdorff d_f (geometry), walk d_w (kinetics), and spectral $d_s = 2d_f/d_w$ (Einstein relation) (Alexander and Orbach, 1982; Rammal and Toulouse, 1983; Havlin and Ben-Avraham, 1987; ben-Avraham and Havlin, 2000; Barlow, 1998)—each in a different row.

Proposition 13 (Substitution table). *For a CRN with polynomial vertex structure placed on a fractal substrate (with d_f , d_w , d_s defined by the heat kernel and graph geometry), the following rows of the CFAC pipeline take the indicated substrate dependence:*

Physical quantity	Mechanism in CFAC	Substitution for d
Universality class / Puiseux p/q	From $\phi(G)$ and branch point of $G = z\phi(G)$; algebraic.	none (substrate-independent)
Density decay, DLR kinetic exponents	Propagator return probability $P(t) \sim t^{-d_s/2}$.	$d \rightarrow d_s$
UV divergences, loop power counting	Loop measure effectively $d^{d_s} q$.	$d \rightarrow d_s$
CNT droplet barrier $\Phi_c(\Delta f)$	$ B(R) \sim R^{d_f}$, $ \partial B(R) \sim R^{d_f-1}$ in the intrinsic metric.	$d \rightarrow d_f$ (subject to caveat below)
Walk statistics, MSD	Anomalous diffusion $\langle x^2 \rangle \sim t^{2/d_w}$.	d_w separately

Row-by-row derivation. Row 1 (universality). The Puiseux exponent p/q at the branch point of $G = z\phi(G)$ is determined by the Newton polygon of the polynomial ϕ alone (§2.5); the substrate does not enter the algebra.

Row 2 (density decay). For diffusion-limited reactions the density obeys $\rho(t)^{-1} \sim t^{d_s/2}$ at long times by Alexander–Orbach–Havlin (Alexander and Orbach, 1982; Havlin and Ben-Avraham, 1987; ben-Avraham and Havlin, 2000); the kernel $P(t, x, x) \sim t^{-d_s/2}$ is the return probability of a simple random walk on the substrate. d_s is the *only* substrate scalar that enters through the propagator in this regime.

Row 3 (UV / loop counting). In the spectral representation, a one-loop integral $\int d^d q f(q^2)/(Dq^2 + M)^\nu$ on a substrate with spectral dimension d_s has an effective integration measure scaling as $d^{d_s} q$ by Hattori–Hattori–Watanabe and Burioni–Cassi (Hattori et al., 1987; Burioni and Cassi, 1996); the engineering dimensions of fields and couplings are read off with $d \rightarrow d_s$ in the power-counting relation of (C4). The upper critical dimension d_c migrates to the value in terms of d_s . This is the ϕ^4 statement quantitatively verified on graphs with tunable d_s by Cescatti et al. 2024 (Cescatti et al., 2024); the ε -expansion is organised around $\varepsilon = d_c - d_s$.

Row 4 (CNT droplet barrier). CNT is a thermodynamic saddle of $\Phi(R) = -\Delta f \cdot |B(R)| + \sigma \cdot |\partial B(R)|$, where $|B(R)|$ is the substrate volume of a ball of radius R and $|\partial B(R)|$ its boundary. For any substrate where the volume obeys $|B(R)| \sim R^{d_f}$ and the boundary obeys $|\partial B(R)| \sim R^{d_f-1}$, extremisation yields $R_c \sim \sigma/\Delta f$ and

$$\Phi_c \sim \frac{\sigma^{d_f}}{\Delta f^{d_f-1}}. \quad (28)$$

The derivation is the standard generalised-dimension CNT calculation (Langer, 1969; Binder and Stauffer, 1976; Binder, 2016) with $d \rightarrow d_f$.

Row 5 (walk statistics). Independent of the above: $\langle x^2(t) \rangle \sim t^{2/d_w}$ is the definition of the walk dimension; d_w enters the Zeldovich/attachment rate in the nucleation prefactor but not the barrier itself. $\square$

Proposition 13 is not a single conjecture but an assembly of standard results, each previously established in its own literature; the clean statement as a row-by-row decomposition matched to the CFAC pipeline is, as far as we can find, not previously in the literature in this form. Rows 1, 2, 3, 5 are thoroughly standard. Row 4 is the apparently-new prediction and requires care.

Critical caveat on Row 4 (finite ramification). The boundary scaling $|\partial B(R)| \sim R^{d_f-1}$ is *not automatic* on every fractal. It holds on infinitely-ramified fractals (Sierpinski carpets, Menger-type, percolation clusters in the scaling limit) but *fails* on finitely-ramified / post-critically-finite fractals (Kigami, 2001): on the Sierpinski gasket, for instance, the boundary of an intrinsic ball is a finite set of junction points, $|\partial B(R)| = O(1)$, not R^{d_f-1} . The derivation of (28) therefore applies to infinitely-ramified fractals; on gaskets the correct CNT calculation gives a different (and smaller) barrier because surface cost does not grow with droplet size. This exact scaling behaviour is the reason the finitely-ramified case gives apparent $d_{\text{eff}} > d_f$ in a small-system CNT test (see §12.9): surface is effectively free, so droplets are anomalously easy to form until they saturate the lattice.

Other caveats, stated cleanly. *Log-periodic corrections* from discrete scale invariance (Akkermans–Dunne–Teplyaev (Akkermans et al., 2009), Saleur–Sornette (Saleur and Sornette, 1996)) decorate all five rows with oscillations in $\log t$ (or $\log R$) at subleading order; they do not affect leading scaling. *Ramification, lacunarity, and connectivity* can make Ising / Potts critical exponents on deterministic fractals depend on more than just d_s (Gefen et al., 1983; Monceau et al., 1998); CFAC’s Row 3 is for Gaussian / perturbative sectors and for the quartic regime quantitatively validated by Cescatti et al. (Cescatti et al., 2024), and does not claim to cover strongly-coupled fixed points on exotic fractals.

Validation of Row 2 (kernel). Sierpinski gasket at level 6 (1095 vertices, exact $d_s = 2 \ln 3 / \ln 5 = 1.365$), return probability $P(t) \sim t^{-d_s/2}$ from 8000 random walks at hopping probability $D \in \{0.1, 0.3, 0.5, 1.0\}$:

D	0.1	0.3	0.5	1.0	exact
d_s	1.386	1.361	1.352	1.388	1.365

All within 2% of the exact value, independent of D as required for a graph-intrinsic property, and consistent with Alexander–Orbach (Alexander and Orbach, 1982). Finite-size effects are significant: at level 4 (123 vertices), apparent $d_s \in \{0.14, 0.49\}$; the random walk must sample several hierarchical levels before the spectral exponent stabilises.

Row 4 prediction and the route to verifying it. Equation (28) predicts that for a CRN with polynomial vertex structure on an infinitely-ramified fractal (carpet-type), the log of the CNT barrier should scale linearly with $\log \Delta f$ with slope $-(d_f - 1)$ —for a Sierpinski *carpet* with $d_f = \log 8 / \log 3 \approx 1.893$, the predicted slope is -0.893 . On a finitely-ramified gasket the slope is *not* $-(d_f - 1)$ by the caveat above, and the $\log B$ vs $\log \Delta f$ fit therefore cannot test (28) cleanly. A decisive test requires a carpet substrate with low (single-site) coupling α , so the critical droplet is small compared to the lattice—we flag this as the cleanest available experimental verification of Row 4 (§12.9).

8 Tropical bridge: the one-loop bubble is algebraically exact

A natural question is whether the *bridge table* of loop integrals (§4) can itself be expressed in closed algebraic form, rather than left as numerical quadratures. The answer at one loop is affirmative, and the relevant technology is *tropical geometry* Borinsky (2020); Borinsky et al. (2025); Arkani-Hamed et al. (2022): the integrand is written in Schwinger form, the first Symanzik polynomial F is viewed as a tropical polynomial on the projective simplex of Schwinger parameters, and its Newton polytope determines the asymptotic structure.

For the one-loop DP self-energy bubble in $d = 2$,

$$I_{\text{DP}}(p; m_1, m_2) = \int_0^1 dt \frac{1}{F(t; p, m_1, m_2)}, \quad F = m_1^2 t^2 + (p^2 + m_1^2 + m_2^2) t(1-t) + m_2^2 (1-t)^2,$$

the Newton polytope is the 1-simplex with vertex monomials at $t \rightarrow \{0, 1\}$ and an edge-interior monomial at $t = 1/2$. Applying the Arkani-Hamed–Hillman–Salvatori “leading-facet plus Hessian” prescription Arkani-Hamed et al. (2022) yields the closed form

$$I_{\text{DP}} = \frac{1}{\sqrt{\lambda_E}} \ln \left[\frac{(\sqrt{\lambda_E} + p^2 + m_1^2 + m_2^2)^2}{4 m_1^2 m_2^2} \right], \quad (29)$$

where the Euclidean Källén-like discriminant is

$$\lambda_E(p^2, m_1^2, m_2^2) = (p^2 + m_1^2 - m_2^2)^2 + 4p^2 m_2^2 = \det(\text{tropical Hessian at saddle}).$$

The prefactor $1/\sqrt{\lambda_E}$ is the Gaussian fluctuation determinant around the tropical saddle point; the logarithm is the sum of Newton-vertex boundary contributions.

Proposition 14 (tropical one-loop exactness). *For any one-loop two-propagator diagram in the coupled DP–MSR theory, the AHHS tropical-saddle + Hessian formula (29) evaluates the Schwinger integral exactly, not asymptotically.*

Proof sketch. The Symanzik polynomial F is homogeneous of degree $L + 1 = 2$ in the Schwinger parameters; on the projective simplex it is a non-degenerate quadratic form whose Hessian is constant. A non-degenerate Gaussian integral against $F^{-(d/2)}$ is evaluated exactly by its saddle plus Hessian prefactor; no higher corrections arise because the Taylor expansion of $\ln F$ around the saddle terminates at second order up to a total derivative. The log term collects the boundary contributions from the two Newton vertices. $\square$ $\square$

Numerical verification across 14 orders of magnitude of kinematic ratio p^2/m^2 and m_1^2/m_2^2 confirms (29) to machine precision. By contrast, the crude single-vertex tropical approximation $I \sim 1/F_{\min}$ is off by factors 10^2 – 10^7 in non-symmetric regimes, and a three-cell Newton-polytope dissection without the Hessian overshoots by 10–70%; the Hessian is essential.

Consequences for the CFAC bridge. (a) The coupling-sector one-loop table is not a numerical dictionary but a closed algebraic expression parameterised by (p^2, m_1^2, m_2^2) . (b) Branch cuts (the Minkowski threshold $s = (m_1 + m_2)^2$) appear as the vanishing locus $\lambda_E = 0$ of the Hessian determinant—directly visible on the Newton polytope. (c) For higher-loop diagrams the Symanzik polynomial has degree $L + 1$ and is no longer quadratic; genuine multi-facet sums are required, for which Borinsky’s tropical Monte Carlo [Borinsky \(2020\)](#); [Borinsky et al. \(2025\)](#) gives an exact integration procedure. At one loop—the setting of Proposition 3 (perturbative exactness in the minimal KS coupling sector)—the tropical form is exact and CFAC is purely algebraic from reaction network through fixed-point exponent.

When one loop is the whole story. For CRNs where the superficial-degree enumeration of (C4) establishes *one-loop exactness of the UV divergence structure*—i.e. $\omega(L) < 0$ for every topology with $L \geq 2$, as we prove in Proposition 3 for the minimal KS coupling—Corollary ?? of the companion theorem (?) closes the *entire* perturbative pipeline in closed algebraic form: reaction network $\rightarrow$ vertex dictionary $\rightarrow$ Lagrange equation $\rightarrow$ Puiseux exponent + branch-gap + tropical bridge, with no numerical quadrature anywhere. This is not a narrow coincidence: the standard coupled-field literature (pair annihilation, pure DP, BARW, Gelimson–Golestanian chemotaxis, Pruessner–Garcia-Millan KS) stops at one loop precisely because the -enumeration of those systems rules out higher-loop UV divergences. Where CFAC genuinely fails to close in symbolic form is the residual class of CRNs for which the enumeration does not yield one-loop exactness (e.g. quartic ψ^4 active-birth sectors with persistent two-loop divergences); there one keeps tropical as an algorithmic rather than symbolic tool via Borinsky [Borinsky \(2020\)](#); [Borinsky et al. \(2025\)](#).

9 Comparison with direct-RG programmes

The direct-RG approach to coupled particle-field theories [Pruessner and Garcia-Millan \(2025\)](#); [Bothe et al. \(2023\)](#); [Gelimson and Golestanian \(2015\)](#) typically proceeds as follows. The coupled action is written explicitly; for active matter and chemotactic systems the field is often integrated out to produce a single-field action with nonlocal (e.g. Coulomb) interaction; Feynman rules are derived; one-loop integrals are computed explicitly; fixed points and exponents are extracted from the resulting beta function.

The AC approach differs in emphasis. Side-by-side:

Task	Direct RG (e.g. Pruessner and Garcia-Millan (2025) ; Gelimson and Golestanian (2015))	CFAC (this work)
Enumerate diagrams	by hand or from Feynman rules, system-specific	from reaction network, automatic
Power counting	from field dimensions at each vertex	$\omega = dL - 2I + n_{\theta}^{\text{tot}}$ on enumerated graphs
Loop integrals	computed per system	once, tabulated (bridge)
Instanton (nucleation rate)	complex-field saddle point, system-specific	branch-gap integral of the DSE polynomial, Theorem 6
All-orders claims	require resummation or explicit higher-loop work	direct from graph enumeration (Prop. 3)
Change the system	redo everything	change the vertex dictionary; rest automatic

A reader familiar with the Pruessner–Garcia-Millan or Gelimson–Golestanian calculations can cross-check this paper as follows. Equation (13) is the one-loop bacterial self-energy computed in the two-field formalism; it reduces to the Coulomb-formalism result of [Gelimson and Golestanian \(2015\)](#) upon integrating out the MSR field at steady state. The beta function (14) has the same sign (stabilising) but our formulation establishes *perturbative exactness* via graph enumeration. Equation (20) gives the same instanton action as the standard x -space WKB result; our derivation expresses it as a contour integral around a DSE branch point, a form not given in [Pruessner and Garcia-Millan \(2025\)](#).

The claim of this paper is that the combination of Proposition 3 (exhaustive enumeration), Theorem 6 (branch-gap instanton), and Theorem 10 (bi-colour structural factorisation) gives a cleaner, system-independent computational route. The direct-RG literature can compute (13) for a specific system; CFAC computes the framework once, then reuses it.

10 Illustrative applications

The case studies below are not quantitative predictions intended to compete with specialised fits ([Meisl et al. \(2014\)](#) for amyloid kinetics, or mesoscale cloud models for ice-crystal initiation). They illustrate the CFAC pipeline end-to-end: reaction network $\rightarrow$ vertex dictionary $\rightarrow$ cubic DSE $\rightarrow$ branch-gap barrier $S \rightarrow$ order-of-magnitude lag time. In both systems the dimensionless parameters are *inferred*, with substantial uncertainties, from published rate constants.

Integrating out the field at steady state gives a cubic $w^+(x) = a + \lambda x^2$, $w^-(x) = \delta x + dx^3$ with $\lambda = g\mu/\kappa$.

Amyloid fibril formation. $A\beta 42$ rate constants ([Cohen et al., 2013](#); [Meisl et al., 2014](#)): $k_+ \sim 3 \times 10^6 \text{ M}^{-1} \text{ s}^{-1}$, $k_2 \sim 10^4 \text{ M}^{-2} \text{ s}^{-1}$, $n_c = n_2 = 2$. The effective coupling $\lambda = k_2 m_0^2 \ell_{\text{avg}} / k_{\text{clear}}$ requires two inferred quantities: a cellular clearance rate $k_{\text{clear}} \sim 10^{-5} \text{ s}^{-1}$ (each of k_{clear} and ℓ_{avg} carries at least one order of magnitude uncertainty). With these inferred values, $\lambda \sim 4$ at μM monomer. Equation (20) gives $S \approx 0.66$ and $\tau \approx 8e^{0.66V}$. At $m_0 = 3 \mu\text{M}$ the classical growth-dominated estimate $t_{\text{lag}} \sim 2/\sqrt{2k_+k_2m_0^3} \approx 0.4 \text{ hr}$ is within a factor ~ 3 of measured 1–3 hr at order-of-magnitude ([Meisl et al., 2014](#)); this is a consistency check, not a new prediction beyond the Knowles-group fits. At $m_0 \sim 1 \text{ nM}$ the branch-gap barrier is rate-limiting, which is consistent with—but does not independently establish—the multi-decade pathology timescale ([Knowles et al., 2009](#)).

Ice crystal formation. Same cubic, λ interpreted as a surface-catalysed (Wegener–Bergeron–Findeisen) nucleation rate at representative values. $S \approx 0.79$, $\tau \approx 11 e^{0.79V}$. Within the model, cooling (increasing λ by 10%) accelerates the onset by $5\times$ at $V = 20$; warming (decreasing λ by 10%) delays it by $9\times$. These sensitivities are qualitatively consistent with observed non-monotonic responses of mixed-phase cloud glaciation to turbulence (Pinsky and Khain, 2002), but should not be read as quantitative atmospheric predictions; a real cloud model needs spatial structure, turbulent advection, and size-distribution kinetics that are outside this $d = 0$ treatment.

Both predictions reduce to the *same polynomial pipeline*: discriminant of the cubic, branch-gap integral, numerical quadrature. Simulation cost at $V = 100$: e^{66} time steps.

11 Rare-events demonstration

To demonstrate concretely that CFAC predicts nucleation rates beyond the reach of direct simulation, we run Gillespie simulations of the coupled system ($a=0.1$, $\lambda=5$, $\delta=4$, $d=1$; barrier $S = 0.456$) in the accessible regime $V \leq 18$ and fit the CFAC prefactor C in $\tau = C e^{VS}$ (Fig. 1). Using $C \approx 5.3$ extrapolated from $V = 18$:

V	τ (CFAC prediction)	feasibility of direct simulation
25	4.7×10^5	possible
50	4.2×10^{10}	expensive
100	3.3×10^{20}	infeasible (beyond any realistic budget)
200	2.1×10^{40}	infeasible
500	5×10^{99}	infeasible

The Gillespie points at $V \leq 18$ fall on the straight line $\log \tau = \log C + VS$ (Fig. 1): simulation at $V = 18$ already requires $\sim 10^5$ time steps per trial. Direct simulation at the biologically or physically relevant $V \sim 10^2$ would require $\sim 10^{22}$ time steps per trial—beyond any accessible budget, yet CFAC delivers τ from polynomial algebra with no additional cost at any V .

The rare-events regime is exactly where direct methods—Gillespie, KMC, lattice Monte Carlo—break down, and exactly where CFAC’s algebraic reformulation becomes essential. Theorem 6 says the instanton action is a *polynomial quantity*; the transfer theorem of analytic combinatorics says the rate it governs is an exact asymptotic of the generating function. Neither statement references system size, so the computation has the same cost at $V = 10$ and $V = 10^6$.

12 Discussion

12.1 What AC buys you that direct RG doesn’t

The decomposition (8) is not bookkeeping. It makes three classes of results accessible that are hard to reach by the direct-RG route.

Exhaustive all-orders enumeration. Proposition 3 states that *no* UV divergence exists beyond one loop for the coupled KS vertices. Establishing such a claim in the direct-RG framework requires either a resummation argument or explicit higher-loop computation; the algebraic formalism of §2.1 reduces it to a finite graph enumeration. The superficial degree of divergence $\omega(\Gamma, d)$ can be read off every 1PI graph; the claim “ $\omega < 0$ for all $L \geq 2$ ” is then a finite check, not a resummation.

Non-perturbative rates from polynomial data. Theorem 6 (branch-gap instanton) rewrites the x -space WKB quasi-potential as a z -space contour integral. The two expressions have equal value but different structure: S_{WKB} requires integrating $\ln(w^-/w^+)$ along a deterministic path; S_{gap} requires only the polynomial coefficients of the DSE, its discriminant, and the difference of




Figure 1: CFAC extrapolation of the mean first-passage time τ across the accessible–inaccessible boundary of direct simulation. Red: Gillespie measurements at $V = 5 \dots 18$. Blue: CFAC prediction $\tau = C e^{VS}$ with $S = 0.456$ from the branch-gap formula and $C \approx 5.3$ fit from the simulation at $V = 18$. The red-shaded region marks where simulating a single trajectory requires more than $\sim 10^{12}$ event computations. The CFAC prediction holds uniformly.

its branches. For systems where the reaction network is known but the deterministic trajectory is not in closed form, the z -space formulation is easier. More importantly, the z -space form exposes the instanton as a feature of the *generating function’s singularity structure*—the same object the transfer theorem (7) uses to sum the perturbation series. Perturbative and non-perturbative physics live at the same branch point.

Structural claims about the coupled DSE. Theorem 10 is a statement about the matrix-tree theorem on bi-coloured graph Laplacians. It establishes the p^{L-1} structural factor for all $L \geq 2$, verified numerically through $L = 4$; the residual form $Q_G(p, f)$ may or may not split further over $\mathbb{Q}$, a topology-dependent arithmetic question (Conjecture 11). The effective consequence is that CFAC’s bi-colour bridge table has an $(L - 1)$ -fold monochromatic piece plus a degree- L residual, generalising the single-colour results of Amarteifio (2019); Bogner and Weinzierl (2010).

12.2 Relation to Pruessner–Garcia-Millan and Gelimson–Golestanian

Pruessner and Garcia-Millan Pruessner and Garcia-Millan (2025) and Bothe *et al.* Bothe *et al.* (2023) establish that the DP (particle) sector must be retained: coarse-grained MSR-type theories produce spurious entropy production. They focus on pure particle systems; the coupled particle-field setting is implicit in their framework but not the object of their calculations. The present paper takes their conclusion as a premise and asks the orthogonal question: given that both sectors must be kept, how can the calculation be organised so that the system-specific part (reaction rates, diffusion constants) is factored from the system-independent part (graph topology, Kirchhoff polynomials)? Our answer is the CFAC decomposition.

Gelimson and Golestanian Gelimson and Golestanian (2015) perform a one-loop RG calculation for a related chemotactic system, integrating out the field to obtain a long-range Coulomb interaction. Their one-loop result is the analogue of (13); the sign of the correction (stabilising) agrees. Our Proposition 3 is stronger: the corresponding beta function is exact to all perturbative orders for the two-vertex theory (12), by direct graph enumeration. This is an instance of a more

general phenomenon: for theories whose vertex content is sufficiently constrained by the reaction network, graph enumeration renders the perturbative series trivial beyond a finite loop order.

12.3 DP vs. MSR through the CFAC lens

A question that recurs throughout the Pruessner–Garcia-Millan programme [Pruessner and Garcia-Millan \(2025\)](#); [Bothe et al. \(2023\)](#) is *when* the Doi–Peliti and Martin–Siggia–Rose formulations agree, and when they diverge. The CFAC perspective organises the answer in four cases (full discussion in Appendix E). For *scaling exponents*, the two agree: the singularity type of the DSE generating function depends only on the polynomial degree of the vertex function, not on whether G counts particles (DP) or density fluctuations (MSR). For *entropy production*, DP is correct and coarse-grained MSR is wrong [Pruessner and Garcia-Millan \(2025\)](#); in CFAC language, entropy production depends on the *full* counting structure, not just the singularity type. For *continuous order parameters* (Ising, ϕ^4 , KPZ), MSR is natural and DP is contrived—such systems lack a natural Fock-space structure, and the $\Delta\phi \rightarrow 0$ limit of a discretised DP reduces to the MSR path integral. For *coupled* systems, neither suffices alone: discrete particles need DP, continuous fields need MSR, and CFAC is the combinatorial bridge.

12.4 Why this matters for physics

Coupled particle-field dynamics is a common setting in modern statistical physics: chemotaxis, active matter, population ecology on spatial substrates, catalytic surfaces, atmospheric micro-physics, neurodegenerative aggregation. In all these settings the quantities of interest—critical densities, nucleation rates, mean first-passage times—are *system-specific* (depending on the particular reaction rates and coupling strengths) but are computed from *system-independent* machinery (Feynman rules, loop integrals, saddle-point analysis). The CFAC decomposition puts the system-specific part in the reaction network and the system-independent part in the bridge table. Changing the system changes the input to the pipeline, not the pipeline.

In particular, the branch-gap formula (20) returns the nucleation barrier S from the polynomial coefficients of the DSE via a single numerical quadrature, with no stochastic simulation of the rare event itself. The pre-exponential factor C in $\tau \approx C e^{VS}$ is either calibrated from small- V simulation (as in Fig. 1) or obtained from the transfer-theorem prefactor in (7); the extrapolation to large V uses only S , so for systems where $VS \gtrsim 30$ (a typical cell compartment or cloud parcel) the CFAC prediction holds even where direct Gillespie is infeasible.

12.5 From illustrations to predictions: what would make the case studies predictive?

The amyloid and cloud case studies of §10 are order-of-magnitude consistency checks, not predictions. The uncertainty budget is dominated by two inferred quantities: the cellular clearance rate k_{clear} and the average fibril length ℓ_{avg} for amyloid; the effective nucleation volume V and the surface catalysis rate for clouds. Making the case studies genuinely predictive requires three steps.

First, *independent measurement of the inferred parameters*. For amyloid: k_{clear} from proteostasis turnover assays (pulse-chase or metabolic labelling in cultured cells), ℓ_{avg} from cryo-EM length histograms. These exist in the literature but with specimen- and condition-dependent variation; a systematic collation and uncertainty quantification would fix the two largest unknowns in $\lambda = k_2 m_0^2 \ell_{\text{avg}} / k_{\text{clear}}$.

Second, *a target observable that the existing Knowles-group fits do not already fit*. The classical $\tau \sim m_0^{-(n_2+1)/2}$ scaling ([Knowles et al., 2009](#)) already determines k_+ , k_n , k_2 from bulk ThT kinetics. A new prediction must be in a regime the Knowles fits don’t cover, or on a quantity they don’t target. Two candidates: (i) the *lag-time distribution* (not just the mean) at single-molecule resolution, which depends on $C e^{VS}$ via its variance structure; (ii) the concentration dependence

of the lag at nanomolar monomer, where the branch-gap barrier crosses over from sub-dominant to rate-limiting—a regime outside the typical μM – mM ThT range.

Third, *spatial extension of Theorem 6 to $d > 0$* . Real fibril formation happens near membranes; real ice crystals nucleate in spatial supersaturation gradients. The branch-gap formula is currently a $d = 0$ identity, so its application to $V = \text{cell compartments}$ is an effective 0D reduction whose validity is not independently checked. A proof or numerical demonstration of Theorem 6 on a lattice would turn the spatial reduction into a controlled approximation.

Of these, the first is a data-collation problem, the second is a collaboration with an experimental group, and the third is a theoretical extension that is in-scope for follow-up work (see “what remains” below).

12.6 Contributions, distilled

It is worth stating plainly what this paper adds and what it does not. The paper does *not* compute a new critical exponent, solve a long-standing open problem in active matter, or out-compute the direct-RG programme on its own terms. What it does is restructure how coupled particle-field calculations are *organised*, in a way that makes several objects algebraic that were previously obtained by case-by-case computation.

Five concrete contributions, in order of structural depth:

1. **The recognition theorem** (companion note): the Doi–Peliti/MSR dressed-propagator generating function of any CRN-in-medium is a Lagrange equation $G = z\phi(G)$ in the sense of Flajolet–Sedgewick. The proof is a three-line rewriting of the Dyson–Schwinger equation; its *content* is that a body of mathematics not previously drawn on in stochastic field theory—the transfer theorem, Puiseux classification of universality, Newton-polygon singularity analysis—transfers wholesale. Everything else in this paper is a consequence of this recognition being applied to coupled-field examples.
2. **The branch-gap contour integral** for the nucleation barrier (Theorem 6) puts perturbative (C1) and non-perturbative (C3) physics on the same singularity of the same generating function. Direct-RG handles the two with distinct technology (Feynman diagrams for perturbation theory, complex-field saddle points for instantons); CFAC unifies them. The moderate-coupling extension via complex-conjugate saddle sum (Prop. 7) pushes this through to the regime where the real algebraic formulation fails.
3. **The matrix-tree pencil theorem** for the two-colour Kirchhoff polynomial (Theorem 8), together with the DP–MSR KS structural factorisation (Theorem 10). $\mathcal{K}(G; p, f) = p^{L-1} Q_G(p, f)$ at every loop order, forced by the perfect-matching structure of the field-coloured subgraph. This is the one place the paper states a theorem about coupled-field diagrams that was previously only conjectural and numerically verified. It is the technical centrepiece.
4. **The substitution table for fractal substrates** (Prop. 13). Different physical exponents inherit different dimensions from a non-Euclidean substrate: universality class is substrate-independent, kinetic exponents use d_s , CNT droplet uses d_f , walk statistics use d_w . The row-by-row decomposition is apparently new in this clean form; Row 4 ($\Phi_c \sim \sigma^{d_f} / \Delta f^{d_f-1}$) is the apparently-novel prediction, modulo an important ramification caveat.
5. **The one-loop tropical bridge** (Prop. 14) gives the one-loop integral in closed algebraic form (AHHS leading-facet+Hessian), exact because the Symanzik polynomial is quadratic at one loop. Combined with the one-loop-exactness criterion (Prop. 4) this makes the entire CFAC pipeline closed-form symbolic end-to-end for the CRNs of physical interest (pair annihilation, pure DP, BARW, chemotaxis, KS).

12.7 What we find most interesting

An honest ranking. The **branch-gap contour integral** is the single result we expect physicists to care about most: a non-perturbative rate from the same polynomial that controls the perturbation series, computed by contour integration around a branch point rather than by complex-field saddle-point. The complex-conjugate extension to moderate coupling makes it robust beyond weak coupling.

The **DP–MSR bi-colour structural theorem** is the result we find most mathematically satisfying: a graph-theoretic theorem that the KS vertex signature $(2, 1)$ forces a perfect matching on the field-coloured subgraph, giving a pencil eigenvalue of multiplicity $L - 1$. The mechanism is elementary once seen but wasn’t seen before; its consequence—that bi-colour bridge integrals decompose with an $(L - 1)$ -fold monochromatic factor—says the two colour sectors decouple structurally at every loop order.

The **substitution table** is the result we think is most likely to propagate into other calculations. Once stated as a row-by-row decomposition—with each row of the CFAC pipeline picking up exactly one substrate dimension—it becomes a template anyone working on RD systems on complex substrates can use. Row 4 ($\text{CNT} \sim d_f$) is the prediction we’d most like to see tested cleanly.

The **one-loop tropical bridge** is the result we think is most under-appreciated. “The one-loop bridge has a closed algebraic form” sounds technical; what it *means* is that for one-loop-exact CRNs (the standard case) the whole CFAC pipeline—from reaction network to beta function—is closed in polynomial algebra, with no numerical quadrature, table lookup, or Monte Carlo anywhere. This sharpens what “CFAC factors counting out of coupled particle-field renormalisation” means.

12.8 Future work

Four concrete directions, in decreasing order of confidence that they are tractable and worthwhile.

A second coupled system: the ant/KPZ pipeline. Everything in this paper has been demonstrated on KS; a second coupled example in a *different universality class* is the natural next step. Our primary candidate is a chemotactic ant colony—foragers laying a volatile pheromone trail (MSR field) while a discrete population (DP) senses its gradient—because in the continuum limit the trail field is described by the Kardar–Parisi–Zhang equation (Kardar et al., 1986) and the system inherits a non-trivial Puiseux exponent $p/q = 1/3$ rather than the mean-field $1/2$ of KS. Specifically, we expect CFAC to reproduce the critical recruitment density $\lambda_c \sim \rho^{1/3}$ (observed in pheromone-trail experiments on *L. niger* and *M. sanguinea*; see Deneubourg & Goss, *J. Theor. Biol.* 1989) from the vertex dictionary alone, exercising three parts of the framework that KS does not: (i) the compact-DP / triple-root Puiseux class in (C2); (ii) a KPZ-type $(\nabla\phi)^2$ vertex, which has $\tau_v = 0$ and therefore sits at the critical edge of Proposition 4’s one-loop-exactness condition; (iii) Row 4 of the substitution table in a system where the nucleation event is collective (colony-level switching) rather than local (ice-crystal initiation), testing whether $\Phi_c \sim \sigma^{d_f}/\Delta f^{d_f-1}$ carries the right interpretation for a sociobiological transition. A full writeup ties the CFAC paper into the larger AC+ programme and gives a second physical system on which the pipeline is demonstrated end-to-end.

A second coupled system from the literature: Toner–Tu. For readers who prefer a target whose one-loop RG is already published, the Toner–Tu theory of polar flocking (Toner and Tu, 1995, 1998; Toner, 2012) is the cleanest choice. It is a DP–MSR coupled action in exactly the sense used here: self-propelled particles with number fluctuations (DP) coupled to a velocity and density field (MSR), with an explicit one-loop calculation culminating in the anomalous dynamical exponent $z = 2$ and the spectacular Mermin–Wagner-evading long-range order in $d = 2$. A CFAC re-derivation would exercise the vertex dictionary in a different way: convective $\psi\nabla\psi$ vertices give $\tau_v > 0$, so the one-loop-exactness criterion (Proposition 4) does *not* apply—Toner–Tu flocking is a

genuine case where higher-loop UV divergences persist and the tropical bridge loses its closed form. It is therefore the right stress test for the scope of the framework. The two-colour Kirchhoff pencil (Theorem 8) still applies; whether the DP–MSR–KS-style structural factor extends to flocking is an open question whose answer would immediately generalise the perfect-matching argument. Active nematic systems (Ramaswamy et al., 2003; Chen and Toner, 2013) generalise this further and have similar published one-loop anchors.

CNT droplet scaling on an infinitely-ramified fractal. The substitution-table prediction $\Phi_c \sim \sigma^{d_f}/\Delta f^{d_f-1}$ is honest for substrates where $|\partial B(R)| \sim R^{d_f-1}$, which holds on Sierpinski carpets but fails on gaskets. A clean verification requires a carpet substrate (say $d_f = \log 8/\log 3 \approx 1.893$, predicted CNT slope -0.893) at weak inter-site coupling. We expect the prediction to hold; a clean numerical test would turn Row 4 from “derivation” into “derivation plus verification”.

Complete Galois/monodromy analysis of the coupled resultant. The moderate-coupling branch-gap extension via complex-conjugate saddles (Prop. 7) empirically finds three pairs at $\alpha = 0.3$ on the 3-site ring; a general proof that the reachable conjugate-pair set is always finite and combinatorially determined from the resultant’s algebraic data would upgrade the proposition to a theorem. This is a question in computational algebraic geometry—the complex Riemann surface of a two-variable polynomial—and is a tractable PhD project.

Matrix-tree beyond one-loop and beyond two colours. Theorem 10 covers DP–MSR KS (bacterial self-energy). Chemical-sector self-energies follow by symmetry with the F/P roles swapped. Mixed-sector n -point functions (one ψ external, one ϕ external) are not covered—the perfect-matching argument relies on both externals being the same colour. The natural extension is a matrix-tree statement for graphs with mixed colour externals, likely involving a second Symanzik polynomial $F(G; p, f)$ in place of $\mathcal{K}$. More speculatively: three-colour generalisations for theories with two fields plus a reservoir (a DP–MSR–MSR system, e.g. bacteria + chemoattractant + oxygen) would give three-variable pencils whose factorisation structure is unknown. The matrix-tree technology extends but the specific “perfect matching forces factorisation” mechanism does not generalise mechanically; new structural observations are needed.

Beyond these four: a functional-analytic extension of the recognition theorem itself (to momentum-dependent vertex couplings, giving a Lagrange equation in function space rather than formal power series) would cover the chemotaxis coupling at full momentum without the zero-momentum projection; higher-loop tropical bridges (beyond one loop, where the Symanzik is no longer quadratic) need Borinsky’s multi-facet technology and remain algorithmic rather than closed-form. These are real but distinctly harder.

12.9 Limitations

The paper does not (a) construct the full spatial instanton in $d > 0$ in literal Theorem 6 form. The real-algebraic formulation extends verbatim only in the decoupled limit (agreement to 2.6% on a 3-site test), under-predicts by $\sim 20\%$ at weak inter-site coupling, and loses its real branch points entirely at moderate coupling. The refined form (Proposition 7) extends to moderate coupling by complex analytic continuation of the resultant with a Coleman–Callan sum over complex-conjugate saddle pairs; this achieves $\sim 3\%$ agreement with Gillespie at $\alpha = 0.3$, but lacks a general proof that the reachable conjugate-pair set is always finite and combinatorially enumerable from the resultant (a full Galois/monodromy question). The paper also does not (b) prove the strong form of Conjecture 11 (full $\mathbb{Q}$ -splitting of the residual Q_G in Theorem 10); an $L = 4$ counter-example (prism-sub with irreducible quartic Q_G) is given, so the strong form does *not* hold; the theorem ($p^{L-1} \mid \mathcal{K}$) is the unconditional statement we establish. (c) extrapolate the one-loop result to $\varepsilon = 1$ quantitatively, where the ε -expansion is known to fail for one-loop approximations Täuber (2014); (d) the prefactor C in $\tau = C e^{VS}$ is closed in the worked Schlögl-I example (§5.2) both

as transfer-theorem amplitude $\mathcal{K}$ and as Kramers prefactor C_∞ ; extending the symbolic Kramers calculation to coupled systems beyond the adiabatic-field reduction is the residual open question; (e) handle CRNs for which the -enumeration of Proposition 3 does *not* yield one-loop exactness, for which the tropical bridge (§8) is no longer symbolic but remains algorithmic via Borinsky’s multi-facet tropical Monte Carlo Borinsky (2020); Borinsky et al. (2025). For the standard coupled-field systems (pair annihilation, pure DP, BARW, Gelimson–Golestanian chemotaxis, Pruessner–Garcia-Millan KS) the -enumeration *does* deliver one-loop exactness, so (e) is not a restriction in practice.

3-site lattice test of the branch-gap extension. For a 3-site coupled ring with site-coupling strength α , we computed the coupled DSE, took the resultant, and measured the branch-gap prediction against Gillespie MFPT:

α	S (branch-gap)	S (Gillespie)	$S_{\text{gap}}/S_{\text{Gillespie}}$
0.0	0.4558	0.4441	1.026
0.1	0.4498	0.5735	0.784
0.3	— (complex)	0.8372	—

At $\alpha = 0$ the resultant factorises $R_{\text{full}} = R_{\text{sym}} \cdot R_{\text{asym}}$, with the symmetric factor recovering the single-site cubic. At $\alpha = 0.1$ the asymmetric factor still has a real branch point but its branch-gap integral misses $\sim 20\%$ of the barrier; at $\alpha = 0.3$ the asymmetric factor has no real branch points at all (all six roots are complex conjugate pairs), yet the MFPT remains exponential with a well-defined slope. The nucleation saddle exists; it has just left the real algebraic variety. This delimits the scope of Theorem 6 concretely.

Reproducibility. Code reproducing all numerical results (branch-gap verification, MFPT, one-loop integral, branch-structure plots, Sierpinski spectral dimension, three-loop Kirchhoff factorisation) is at `paper/reproduce.py`; figures in `paper/figures/`.

Acknowledgements. Discussions with G. Pruessner and R. Garcia-Millan on the coupled DP–MSR formalism and its entropy-production subtleties were instrumental in framing the problem. The algebraic reaction-diffusion formalism of §2.1 was developed under G. Pruessner’s supervision in the thesis work Amarteifio (2019).

Appendix A: one-loop integral

Bacterial self-energy at one loop is a two-colour bubble,

$$\Sigma_\psi^{(1)}(k, 0) = \chi\mu \int \frac{d^d q}{(2\pi)^d} \frac{(\mathbf{k} - \mathbf{q}) \cdot \mathbf{q}}{D_A q^2 + m + D_c |\mathbf{k} - \mathbf{q}|^2 + \kappa}. \quad (30)$$

Expanding to $O(k^2)$ and collecting terms ($D = D_A + D_c$, $M = m + \kappa$):

$$\delta D_A^{(1)} = 2\chi\mu D_c I_1 - 2\chi\mu D_c^2 I_2, \quad I_j = \int \frac{d^d q}{(2\pi)^d} \frac{q^{2j}}{(Dq^2 + M)^{j+1}}. \quad (31)$$

At $d = 2 - \varepsilon$: $I_1 = \ln(D\Lambda^2/M)/[4\pi D^2]$, $I_2 = \ln(D\Lambda^2/M)/[4\pi D^3]$ at leading order. Combining, $\delta D_A^{(1)} = [\chi\mu D_c D_A/(\pi D^3)] \ln \Lambda$. Vertex corrections: triangles with $\omega = 0$ naively, but momentum routing gives $\sim q^{-3}$ at large q (Appendix C); $Z_\chi = Z_\mu = 1$.

Appendix B: proof of Theorem 6

Let $\varphi(x) = \int_{x_1}^x \ln[w^-(y)/w^+(y)] dy$, so $S_{\text{WKB}} = \varphi(x_2)$. Set $z(y) = w^-(y)/w^+(y)$. Both x_1, x_2 are fixed points of $\dot{x} = w^+ - w^-$, so $z(x_1) = z(x_2) = 1$. For $z \in (1, z_*)$ there are two real branches $x_{\text{stab}}(z) < x_{\text{unst}}(z)$, merging at the branch point z_* of (19).

Change variables: $d\varphi = (\ln z) dy = (\ln z) x'(z) dz$. Along the stable branch, $\varphi(x_2)^{\text{stab}} = \int_1^{z_*} (\ln z) x'_{\text{stab}}(z) dz$; along the unstable branch, $\varphi(x_2)^{\text{unst}} = \int_{z_*}^1 (\ln z) x'_{\text{unst}}(z) dz$. Integration by parts $\int (\ln z) x'(z) dz = [x(z) \ln z] - \int x(z)/z dz$: boundary terms vanish at $z = 1$ and cancel at z_* (equal branch values). The total is (20).

Numerical verification, seven parameter sets.

a	λ	δ	d	z_*	S_{WKB}	$S_{\text{gap}}/S_{\text{WKB}}$
0.50	3.0	2.5	1.0	1.100	0.039071	1.000000
0.30	3.0	2.0	1.0	0.977	0.029733	1.000000
0.30	3.5	2.0	1.0	1.089	0.002200	1.000000
0.10	3.5	3.0	1.0	1.248	0.494986	1.000000
0.02	5.5	5.0	1.0	7.543	0.791624	1.000000
0.30	2.5	2.0	0.5	1.092	0.060758	1.000000
0.05	6.0	4.0	2.0	1.174	0.475059	1.000000

Appendix C: higher-loop enumeration

For 3-leg vertices and $E = 2$ external legs, $3V = 2I + E \Rightarrow V = 2L, I = 2L + 1$. At $L = 2$: the unique 1PI topology is $K_4 \setminus e$ (six vertices, eight edges is $L = 3$; at $L = 2$ the topology has four vertices, five edges). Vertex content $\chi^2 \mu^2$; propagators $3G_\psi + 2G_\phi$; $n_\partial^{\text{tot}} = 4$; $\omega = -2$. At $L = 3$: vertex content $\chi^3 \mu^3$; propagators $5G_\psi + 3G_\phi$; $n_\partial^{\text{tot}} = 6$; $\omega = -4$. The branching vertex $\tilde{\psi}^2 \psi$ and MSR noise vertex $\tilde{\phi}^2$, if added, give no new divergences at these orders either.

Appendix D: three-loop Kirchhoff computation

Three-loop topology: $K_4 \setminus e$, vertices $\{0, \dots, 5\}$ (external at 0, 5), edges $\{(0, 1), (1, 2), (2, 5), (0, 3), (3, 4), (4, 5), (1, 4)\}$. Weighted Laplacian with edge weights $u_1, \dots, u_8$; principal-minor determinant gives $\mathcal{K}$ (36 monomials).

For each of $2^8 = 256$ substitutions $u_i \in \{p, f\}$, applying `sympy.factor`: 254 factorise non-trivially; two exceptions are monochromatic p^8, f^8 . Restricting to physical colourings $(3p + 5f), (4p + 4f), (5p + 3f)$ ($\binom{8}{k}$ for $k = 3, 4, 5$, total 182): all factorise.

Appendix E: DP vs. MSR—four cases through CFAC

The CFAC perspective classifies the relationship between Doi–Peliti (DP) and Martin–Siggia–Rose (MSR) formulations into four cases, organised by *what observable* is being computed. The key quantities in CFAC are: the singularity *type* of the DSE generating function (sets the scaling exponents), the singularity *location* (sets the critical point), the full *coefficient sequence* $[z^n]G(z)$ (sets the distribution of microstates), and the *counting variable* itself (sets what is labelled and distinguishable).

Case 1: scaling exponents—DP and MSR agree. Near a continuous phase transition, exponents depend only on symmetries, dimensionality, and conservation laws. In CFAC language, the singularity type of the DSE kernel $\phi(G)$ (square-root, cube-root, pole) is determined by the polynomial degree and discriminant structure of ϕ , which depends on the vertex *types*, not on whether G counts particles (DP) or density fluctuations (MSR). Both formulations therefore give the same Puiseux exponent and the same mean-field exponents.

Case 2: entropy production—DP correct, MSR wrong. Pruessner and Garcia-Millan [Pruessner and Garcia-Millan \(2025\)](#) established that coarse-grained MSR theories produce spurious entropy production because they erase the microscopic stochastic trajectory information carried by particle identity. In CFAC language, entropy production depends on the *full counting structure*, not just the location of the dominant singularity: it involves (at pair interactions) the three-point density correlator, which in turn depends on the coefficient sequence $[z^n]G$, not just z_* . When MSR “forgets” the integer labelling n and passes to a continuous density, it loses the fine-grained microstate count, and the entropy comes out wrong ([Bothe et al., 2023](#)).

Case 3: continuous order parameters—MSR natural, DP contrived. For systems described by a continuous order parameter with no underlying particle interpretation—Ising magnetisation, ϕ^4 field, KPZ height—there is no natural Fock-space structure to start from. Writing a DP theory requires discretising the field with an artificial lattice spacing $\Delta\phi$ that must later be sent to zero. In CFAC language, the discrete generating function has coefficients indexed by $n = \phi/\Delta\phi$, and the $\Delta\phi \rightarrow 0$ limit transforms the discrete generating function into a continuous Laplace or Fourier transform. This suggests the natural domain of MSR is problems where $\Delta\phi \rightarrow 0$ can be taken *before* the combinatorics, while DP’s natural domain is problems where the integer structure $n \in \mathbb{Z}_{\geq 0}$ survives and matters.

Case 4: coupled particle-field systems—CFAC territory. This is the present paper’s subject. Neither formulation alone suffices: discrete particles need DP (absorbing state, correct entropy production, discrete nucleation); continuous fields need MSR (gradient energies, spatial correlations, continuous order parameters). CFAC organises the coupled calculation by giving both sectors a common generating-function-based footing: the particle sector contributes a DP-style DSE, the field sector contributes an MSR-style propagator, and the coupling vertices produce mixed diagrams with two-colour Kirchhoff polynomials. Theorem 6 (the branch-gap formula) and Proposition 3 (one-loop exactness) are instances of CFAC delivering results that neither DP nor MSR alone makes accessible.

The four cases taken together suggest that *every* stochastic field theory sits somewhere on a spectrum defined by how much of the integer counting structure survives at large V . Case 1 says the spectrum doesn’t matter at criticality. Cases 2–4 say it does everywhere else.

References

- Alexander, S. and Orbach, R., 1982. Density of states on fractals: “fractons”. *J. Physique Lett.* 43, L625.
- Amarteifio, S., 2019. *Field-theoretic formulation and empirical tracking of spatial processes*. PhD dissertation (supervised by G. Pruessner), Imperial College London.
- Amarteifio, S., 2026. *BRW via Analytic Combinatorics: Complete Worked Example*. Companion note, Percolation Labs.
- Bogner, C. and Weinzierl, S., 2010. Feynman graph polynomials. *Int. J. Mod. Phys. A* 25, 2585. arXiv:1002.3458.
- Bothe, M., Cocconi, L., Zhen, Z., and Pruessner, G., 2023. Particle entity in the Doi–Peliti and response field formalisms. *J. Phys. A* 56, 175002. arXiv:2205.10409.
- Cohen, S. I. A. et al., 2013. Proliferation of amyloid- β 42 aggregates occurs through a secondary nucleation mechanism. *PNAS* 110, 9758.
- De Dominicis, C., 1976. Techniques de renormalisation de la théorie des champs et dynamique des phénomènes critiques. *J. Physique Colloq.* 37, C1-247.

- Doi, M., 1976. Second quantization representation for classical many-particle systems. *J. Phys. A* 9, 1465.
- Dykman, M. I., Mori, E., Ross, J., and Hunt, P. M., 1994. Large fluctuations and optimal paths in chemical kinetics. *J. Chem. Phys.* 100, 5735.
- Elgart, V. and Kamenev, A., 2004. Rare event statistics in reaction-diffusion systems. *Phys. Rev. E* 70, 041106.
- Flajolet, P. and Sedgewick, R., 2009. *Analytic Combinatorics*. Cambridge Univ. Press.
- Freidlin, M. and Wentzell, A., 2012. *Random Perturbations of Dynamical Systems*. Springer.
- Gelimson, A. and Golestanian, R., 2015. Collective dynamics of dividing chemotactic cells. *Phys. Rev. Lett.* 114, 028101.
- Arkani-Hamed, N., Hillman, A., and Salvatori, S., 2022. Tropical amplitudes and the method of regions. *arXiv:2202.12296*.
- Borinsky, M., 2020. Tropical Monte Carlo quadrature for Feynman integrals. *arXiv:2008.12310*. [Published 2023, *Ann. Inst. H. Poincaré D* 10, 635.]
- Borinsky, M., Munch, H.J., and Tellander, F., 2025. Tropical Feynman integration for massive scalar theories. *J. High Energy Phys.* (to appear).
- Hattori, K., Hattori, T., and Watanabe, H., 1987. Gaussian field theories on general networks and the spectral dimensions. *Prog. Theor. Phys. Suppl.* 92, 108.
- Havlin, S. and Ben-Avraham, D., 1987. Diffusion in disordered media. *Adv. Phys.* 36, 695.
- Rammal, R. and Toulouse, G., 1983. Random walks on fractal structures and percolation clusters. *J. Physique Lett.* 44, L13.
- ben-Avraham, D. and Havlin, S., 2000. *Diffusion and Reactions in Fractals and Disordered Systems*. Cambridge Univ. Press.
- Barlow, M. T., 1998. Diffusions on fractals. In: *Lectures on Probability Theory and Statistics* (Saint-Flour 1995). Lecture Notes in Math. 1690, Springer, pp. 1–121.
- Burioni, R. and Cassi, D., 1996. Universal properties of spectral dimension. *Phys. Rev. Lett.* 76, 1091.
- Cescatti, F., Ibáñez-Berganza, M., Vezzani, A., and Burioni, R., 2024. Universal scaling in real dimension. *Nat. Commun.* 15, 4236.
- Kigami, J., 2001. *Analysis on Fractals*. Cambridge Univ. Press.
- Akkermans, E., Dunne, G. V., and Teplyaev, A., 2009. Physical consequences of complex dimensions of fractals. *Europhys. Lett.* 88, 40007.
- Saleur, H. and Sornette, D., 1996. Complex exponents and log-periodic corrections in frustrated systems. *J. Phys. I France* 6, 327.
- Gefen, Y., Aharony, A., and Mandelbrot, B. B., 1983. Phase transitions on fractals, I & II. *J. Phys. A* 16, 1267.
- Monceau, P., Perreau, M., and Hébert, F., 1998. Magnetic critical behavior of the Ising model on fractal structures. *Phys. Rev. B* 58, 6386.

- Binder, K. and Stauffer, D., 1976. Statistical theory of nucleation, condensation and coagulation. *Adv. Phys.* 25, 343.
- Binder, K., 2016. Overview: understanding nucleation phenomena from simulations of lattice gas models. *J. Chem. Phys.* 145, 211707.
- Langer, J. S., 1969. Statistical theory of the decay of metastable states. *Ann. Phys.* 54, 258.
- Gardiner, C., 2009. *Stochastic Methods: A Handbook for the Natural and Social Sciences*, 4th ed., Springer.
- Golub, G. H. and Van Loan, C. F., 2013. *Matrix Computations*, 4th ed., Johns Hopkins Univ. Press.
- Gaveau, B., Moreau, M., and Toth, J., 1996. Decay of the metastable state: different predictions between discrete and continuous descriptions. *J. Stat. Phys.* 85, 419.
- Kardar, M., Parisi, G., and Zhang, Y.-C., 1986. Dynamic scaling of growing interfaces. *Phys. Rev. Lett.* 56, 889.
- Toner, J. and Tu, Y., 1995. Long-range order in a two-dimensional dynamical XY model: how birds fly together. *Phys. Rev. Lett.* 75, 4326.
- Toner, J. and Tu, Y., 1998. Flocks, herds, and schools: a quantitative theory of flocking. *Phys. Rev. E* 58, 4828.
- Toner, J., 2012. Reanalysis of the hydrodynamic theory of fluid, polar-ordered flocks. *Phys. Rev. E* 86, 031918.
- Ramaswamy, S., Simha, R. A., and Toner, J., 2003. Active nematics on a substrate: giant number fluctuations and long-time tails. *Europhys. Lett.* 62, 196.
- Chen, L. and Toner, J., 2013. Universality for moving stripes: a hydrodynamic theory of polar active smectics. *Phys. Rev. Lett.* 111, 088701.
- Janssen, H. K., 1976. On a Lagrangean for classical field dynamics and renormalization group calculations of dynamical critical properties. *Z. Phys. B* 23, 377.
- Weinberg, S., 1960. High-energy behavior in quantum field theory. *Phys. Rev.* 118, 838.
- Zimmermann, W., 1969. Convergence of Bogoliubov's method of renormalization in momentum space. *Commun. Math. Phys.* 15, 208.
- Keller, E. F. and Segel, L. A., 1970. Initiation of slime mold aggregation viewed as an instability. *J. Theor. Biol.* 26, 399.
- Knowles, T. P. J. et al., 2009. An analytical solution to the kinetics of breakable filament assembly. *Science* 326, 1533.
- Coleman, S., 1977. Fate of the false vacuum: semiclassical theory. *Phys. Rev. D* 15, 2929.
- Callan, C. G. and Coleman, S., 1977. Fate of the false vacuum II: first quantum corrections. *Phys. Rev. D* 16, 1762.
- Langer, J. S., 1967. Theory of the condensation point. *Ann. Phys.* 41, 108; 1969. Statistical theory of the decay of metastable states. *Ann. Phys.* 54, 258.
- Martin, P. C., Siggia, E. D., and Rose, H. A., 1973. Statistical dynamics of classical systems. *Phys. Rev. A* 8, 423.

- Meisl, G. et al., 2014. Differences in nucleation behavior underlie the contrasting aggregation kinetics of the A β 40 and A β 42 peptides. *PNAS* 111, 9384.
- Peliti, L., 1985. Path integral approach to birth-death processes on a lattice. *J. Physique* 46, 1469.
- Pinsky, M. and Khain, A. P., 2002. Effects of in-cloud nucleation and turbulence on droplet spectrum formation in cumulus clouds. *Q. J. R. Meteorol. Soc.* 128, 501.
- Pruessner, G. and Garcia-Millan, R., 2025. Field theories of active particle systems and their entropy production. *Rep. Prog. Phys.* arXiv:2211.11906.
- Schlögl, F., 1972. Chemical reaction models for non-equilibrium phase transitions. *Z. Physik* 253, 147.
- Täuber, U. C., 2014. *Critical Dynamics: A Field Theory Approach to Equilibrium and Non-Equilibrium Scaling Behavior*. Cambridge Univ. Press.
- Assaf, M. and Meerson, B., 2017. WKB theory of large deviations in stochastic populations. *J. Phys. A* 50, 263001.
- van Kampen, N. G., 1992. *Stochastic Processes in Physics and Chemistry*. North-Holland.